





Chapter 4

Threads

Part **A**: Section 4.1




Motivational Example: *Quicksort*



```
Recu_proc Quicks(Arr, lo, hi)
{
  if ( lo < hi ) then
  {
    p := call Parti(Arr, lo, hi);
    call Quicks(Arr, lo, (p-1));
    call Quicks(Arr, (p+1), hi);
  }
}
```

Those students, who have not yet seen the quicksort algorithm before, are requested to study it in their own time from the CS literature.



- **Remarks:**
- The **Partitioning** procedure does not only yield the split-value **p**, but also has a (pre-sorting) *side-effect* on the **Array**
- Since the two recursive calls of **Quicksort** work on *disjoint* parts of the **Array**, they can at least in principle be carried out **in parallel (concurrently)**, *if* the memory, which holds the **Array**, is allowed to be “shared” by the two recursive invocations.
- By parallelisation **we may hope to save some run-time, if the coordination of parallelism** (“*management-overhead*”) **does not “eat” the time that was saved by parallelisation**

Overview



- **Processes and threads**

- Multithreading
- Thread functionality

- **Types of threads**

- User level and kernel level threads

- **Multicore and multithreading**

- Performance of Software on Multicore

- *Windows process and thread management*

- Management of background tasks and application lifecycles
- Windows process
- Process and thread objects
- Multithreading
- Thread states
- Support for OS subsystems

- *Solaris thread and SMP management*

- Multithreaded architecture
- Motivation
- Process structure
- Thread execution
- Interrupts as threads

- *Linux process and thread management*

- Tasks/threads/namespaces

- *Android process and thread management*

- Android applications
- Activities
- Processes and threads

- *Mac OS X grand central dispatch*

Processes and Threads



Resource Ownership

Process includes a virtual address space to hold the process image

- The OS protects processes and resources to prevent unwanted interference



Scheduling/Execution

Follows an execution path that *may be interleaved with other processes*

- A process has an execution state (Running, Ready, etc.) and a dispatching priority, and is the entity that is scheduled and dispatched by the OS



Processes and *Threads*

- The **unit of dispatching** is referred to as a *thread* or *lightweight process*
- The **unit of resource ownership** is referred to as a *process* or *task*
- **Multithreading** - The ability of an OS to support multiple, concurrent paths of execution within a single process



Motivational Example: *Concurrent Quicksort*

```
Recu_proc Quicks(Arr, lo, hi)
{
  if ( lo < hi ) then
  {
    p := call Parti(Arr, lo, hi);
    call Quicks(Arr, lo, (p-1));
    call Quicks(Arr, (p+1), hi);
  }
}
```

In our concurrent quicksort example:

- A **Process** would be the *owner* of the **Array**, as well as the *owner* of this **program code** (algorithm)
- **Threads**, which can be independently *scheduled* (although they also “belong” to that Process), **would do the work on the fields of the **Array**.**
- Thereby every (recursive) call would “*spawn*” a new thread into existence





Single Threaded Approaches

- A single thread of execution per process, in which the concept of a thread is not recognized, is referred to as a single-threaded approach
- **MS-DOS** is an example




Figure 4.1 Threads and Processes

Multithreaded Approaches

- The right half of Figure 4.1 depicts multithreaded approaches

- A Java run-time environment is an example of a system of one process with multiple threads




Figure 4.1 Threads and Processes

Process in a multi-threaded environment

- Defined as “the **unit of resource allocation** and a **unit of protection**”
- **Associated** with processes:
 - A **virtual address space** that holds the **process image**
 - **Protected access** to:
 - Processors (CPUs)
 - Other processes (for interprocess communication)
 - Files
 - I/O resources (devices and channels)





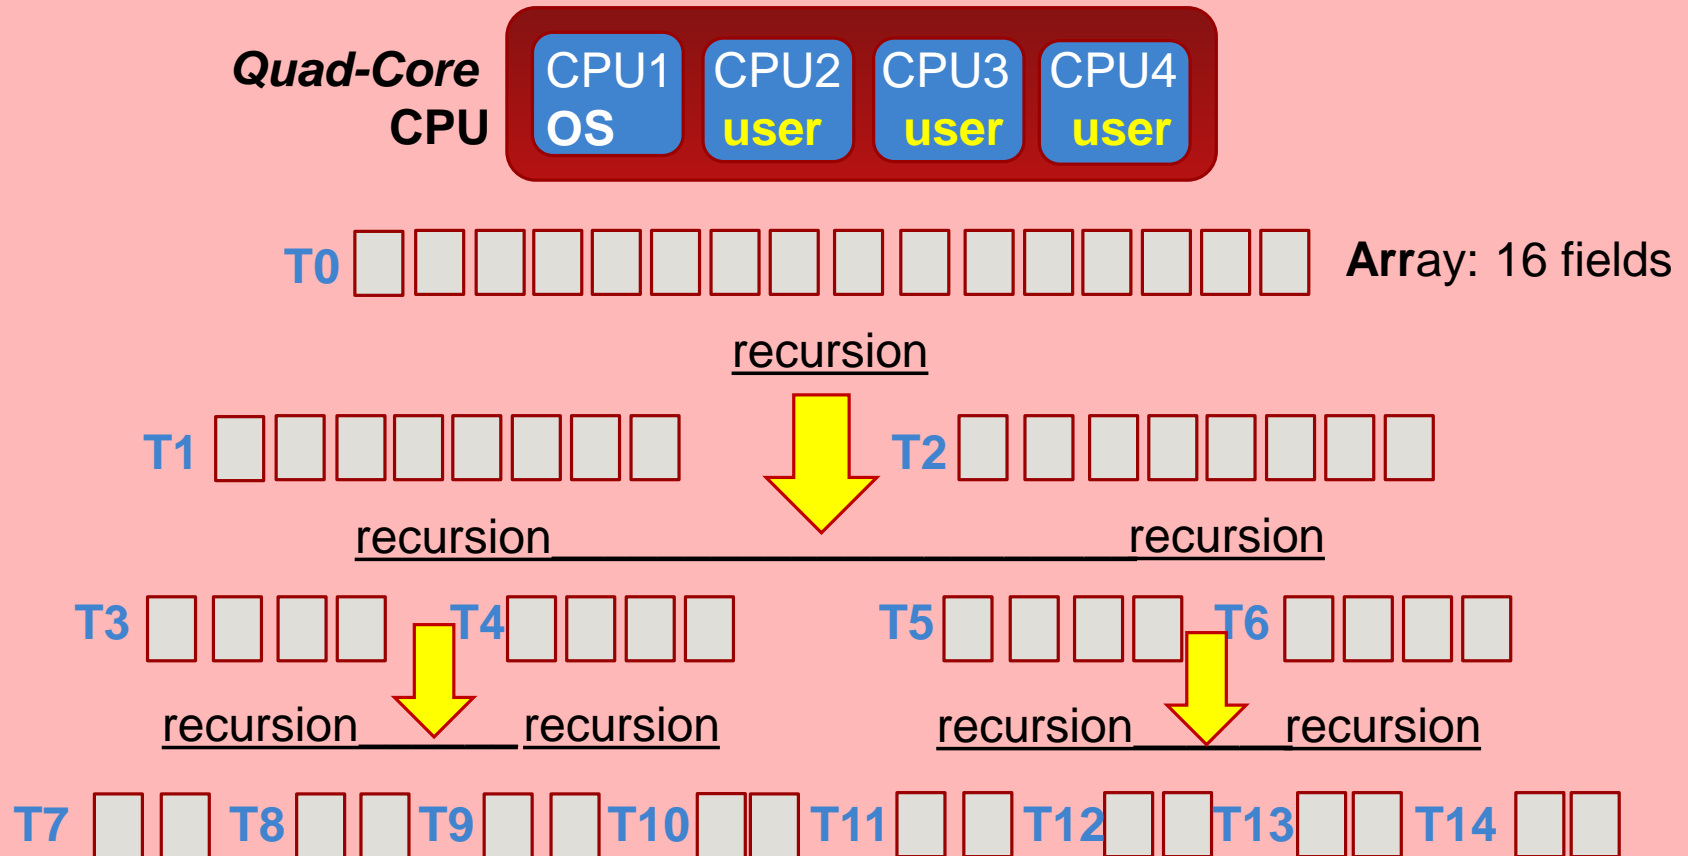
One or More *Threads* in a Process

Each thread of a process has:

- An execution **state** (Running, Ready, etc.)
- A saved thread context when not running (Thread-Control-Block **TCB**)
- An execution **stack** (run-time stack)
- Some storage for its own local variables
- **Access to the memory and resources of its processes**, shared with all other threads in that process




Our Example again: *Concurrent Threaded Quicksort*



After only a few recursions, we have already “spawned” far more threads (T...) than we have CPUs, and from that time onwards the thread-management efforts (scheduling, etc.) begin to “eat” the advantages of parallel computation !



Single-Threaded Process Model



Multithreaded Process Model

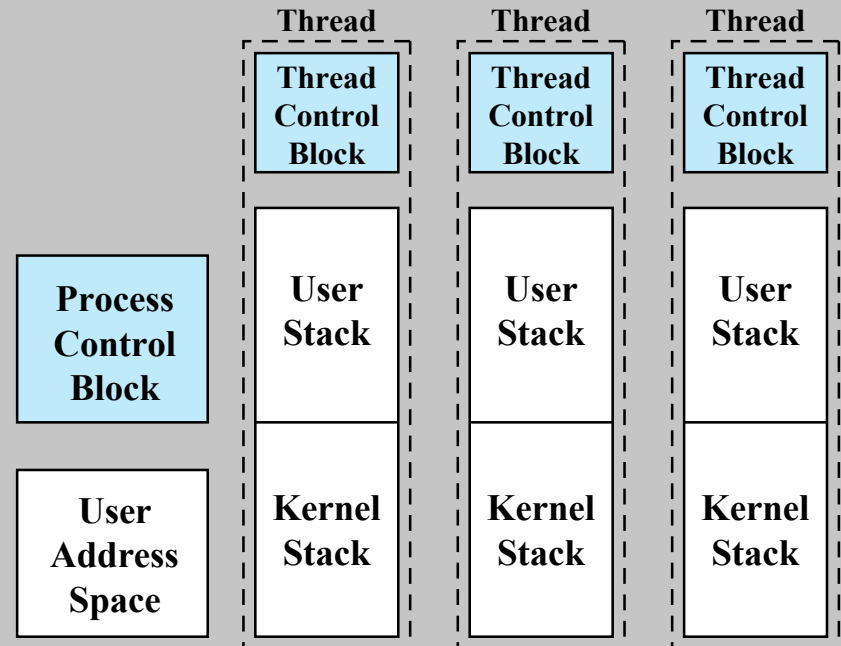
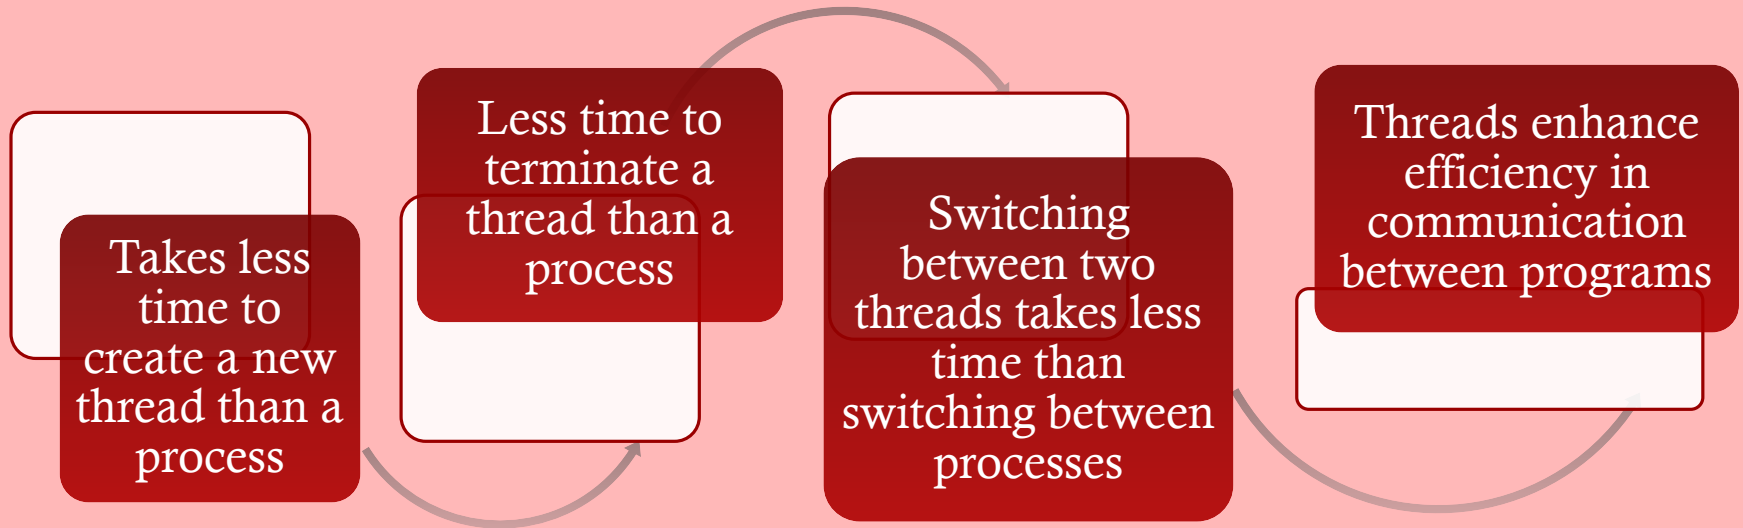


Figure 4.2 Single Threaded and Multithreaded Process Models



Benefits of Threads



Thread Management



- In an OS that supports threads, scheduling and dispatching is done on a thread basis





Most of the state information dealing with execution is maintained in thread-level data structures

- Suspending a process involves suspending all threads of the process
- Termination of a process terminates all threads within the process



Thread Execution States



The key states for a thread are:

- Running
- Ready
- Blocked

Thread operations associated with a change in thread state are:

- Spawn
- Block
- Unblock
- Finish

Example: Benefits of Threading

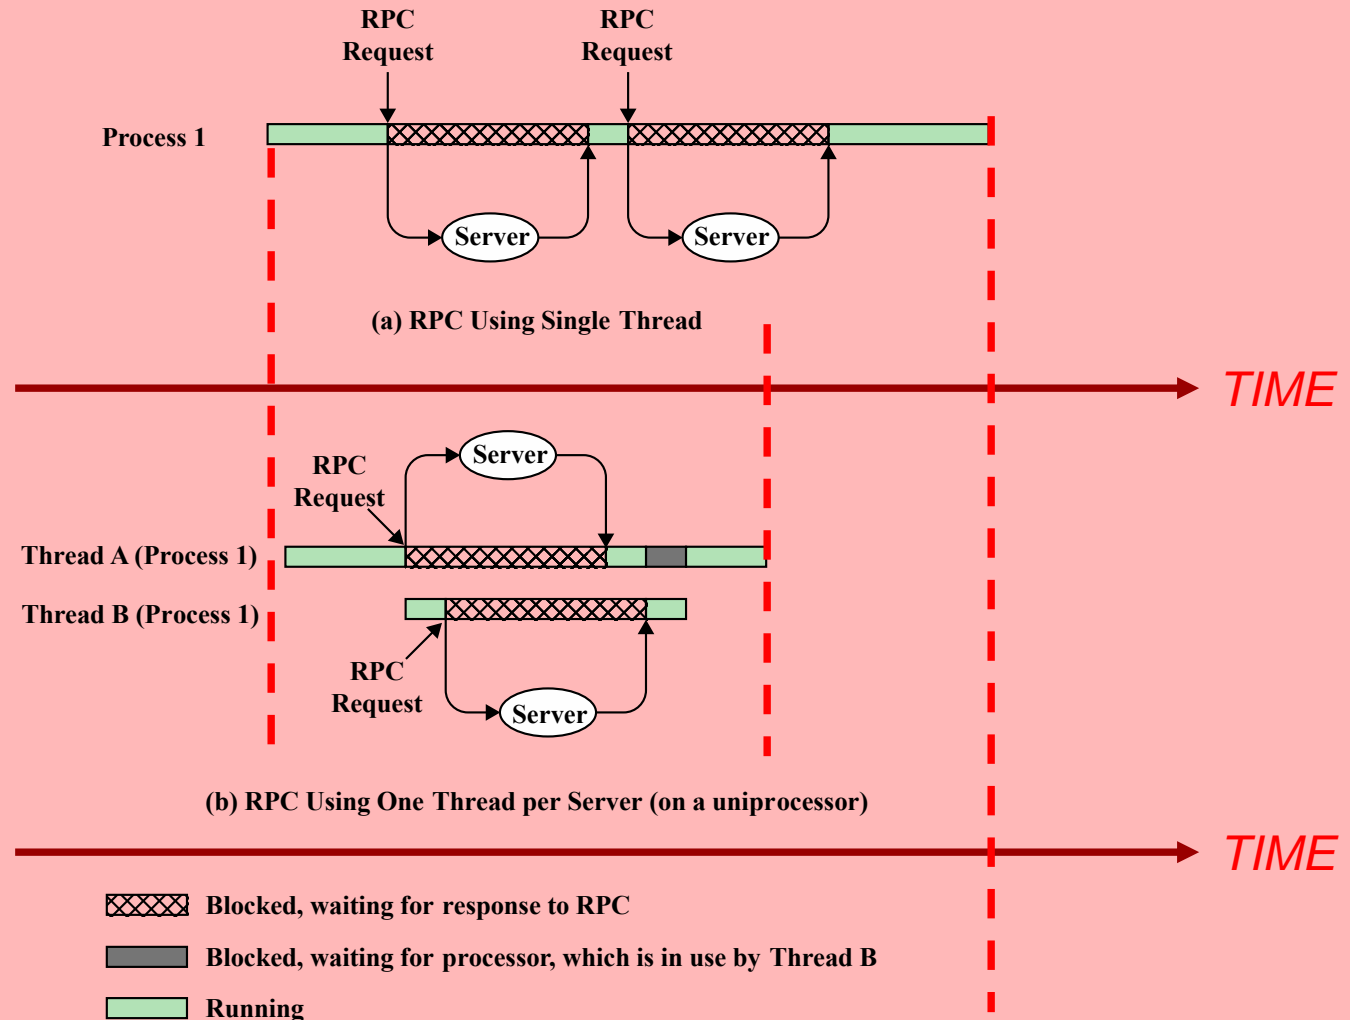


Figure 4.3 Remote Procedure Call (RPC) Using Threads

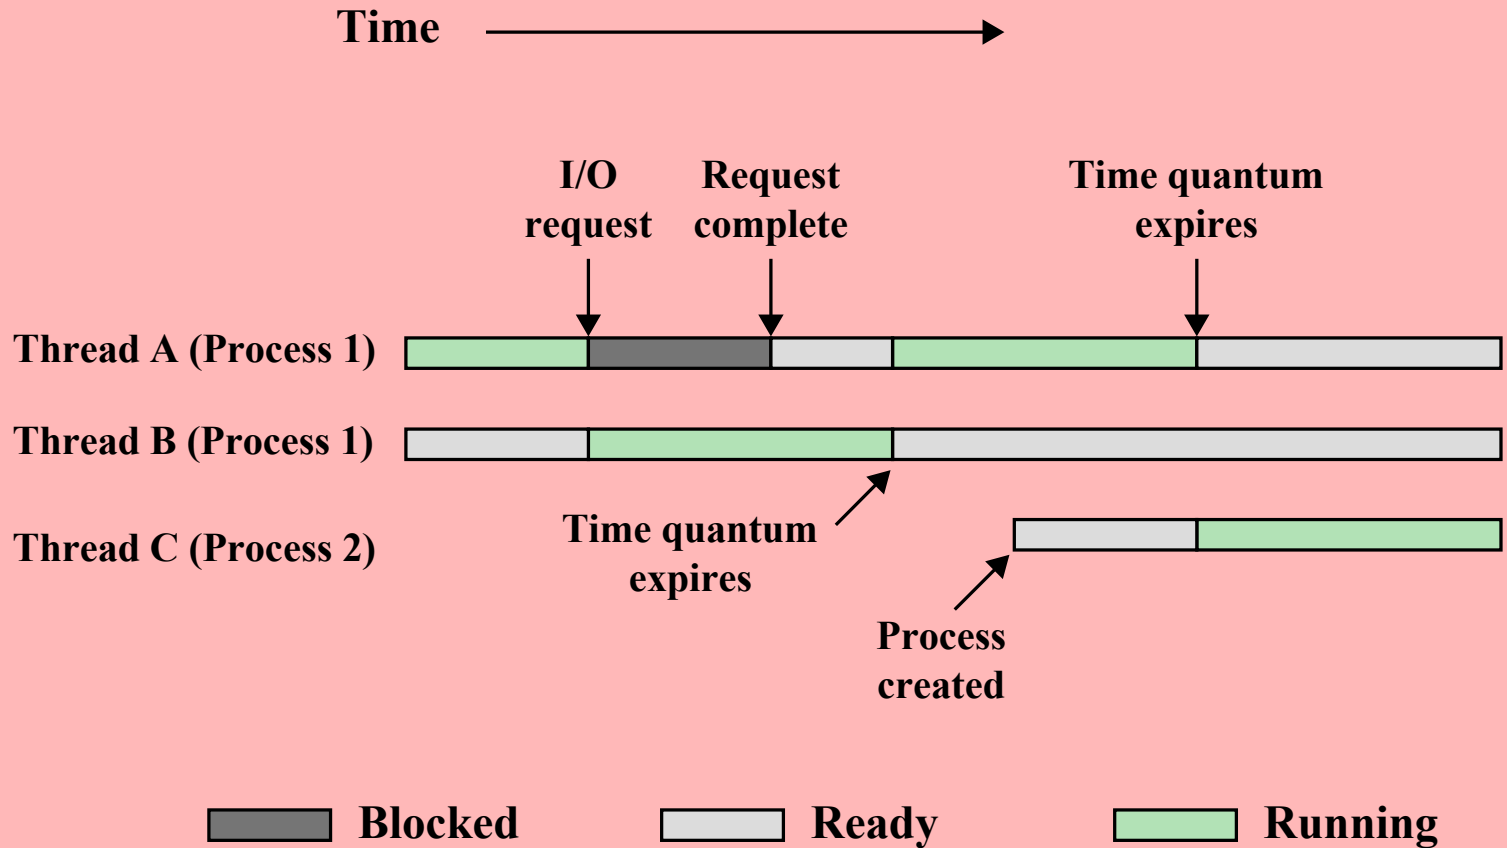


Figure 4.4 Multithreading Example on a Uniprocessor

Thread Synchronization



- It is necessary to synchronize the activities of the various threads
 - All threads of a process share the same address space and other resources
 - Any alteration of a resource by one thread affects the other threads in the same process





Types of Threads

The diagram consists of two large, dark red arrows pointing in opposite directions. The left arrow points left and contains the text 'User Level Thread (ULT)'. The right arrow points right and contains the text 'Kernel level Thread (KLT)'. The two arrows are connected at their bases, forming a continuous shape.

User Level
Thread (ULT)

Kernel level
Thread (KLT)


User-Level Threads (ULTs)




- All thread management is done by the application



- The kernel is not aware of the existence of threads



(a) Pure user-level

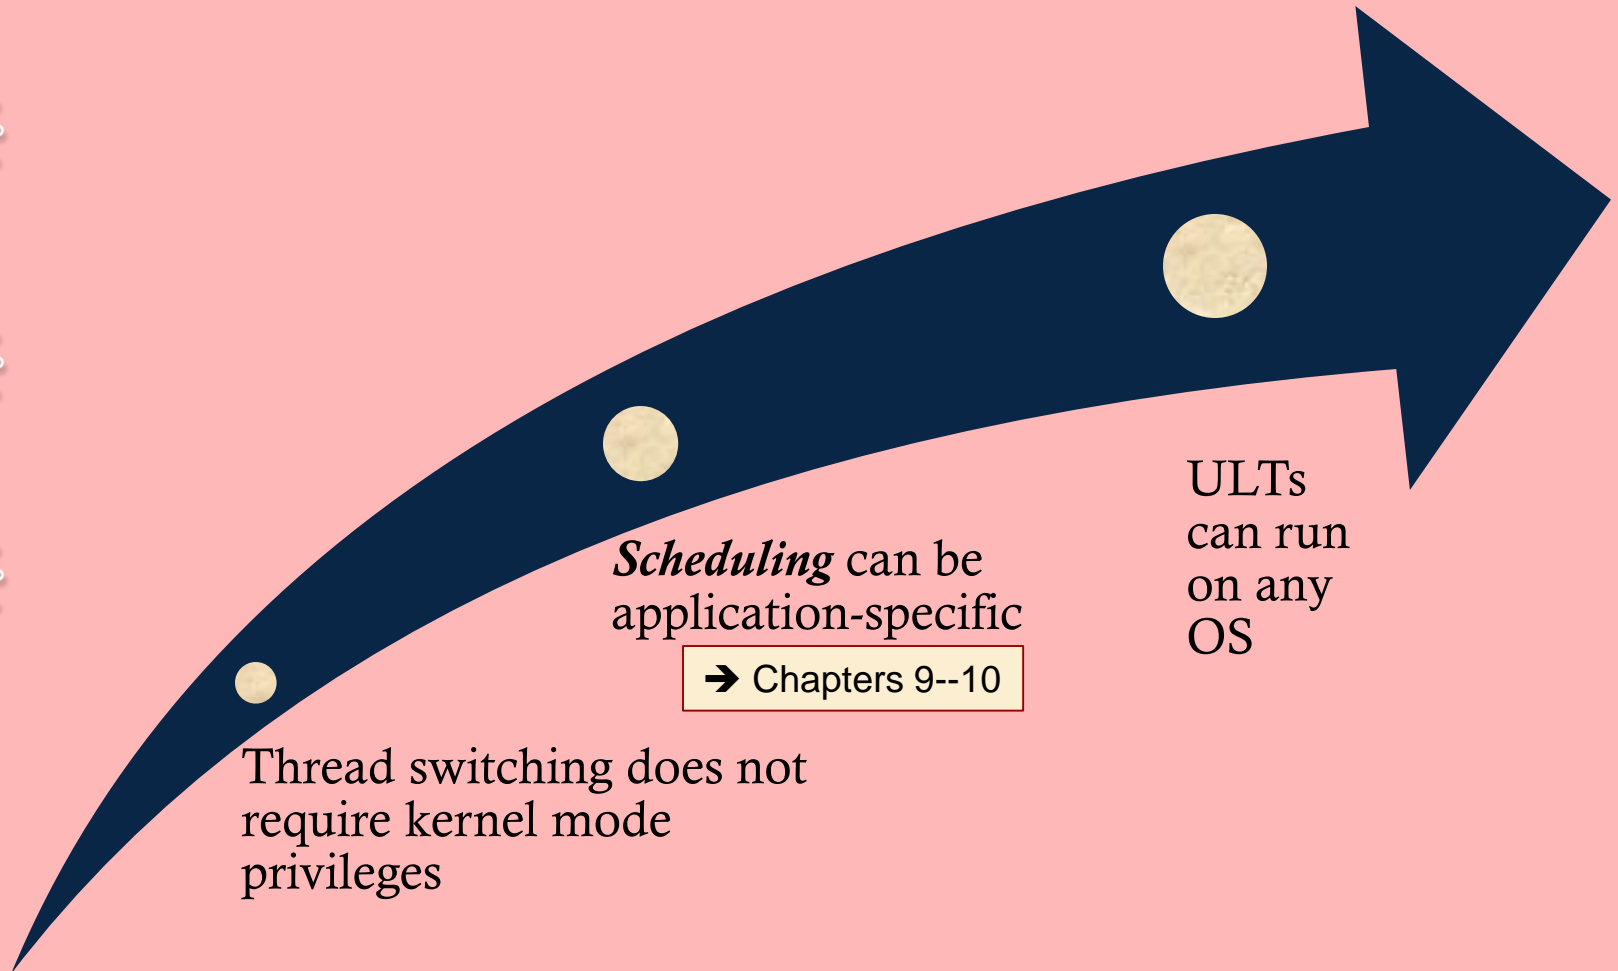


Colored state
is current state

Figure 4.6 Examples of the Relationships Between User-Level Thread States and Process States



Advantages of ULTs



Disadvantages of ULTs

- In a typical OS many system calls are blocking




- As a result, when a ULT executes a system call, not only is that thread blocked, but all of the threads within the process are blocked as well

- In a pure ULT strategy, a multithreaded application cannot take advantage of multiprocessing



- A kernel assigns one process to only one processor at a time, therefore, only a single thread within a process can execute at a time

Old proverb in computer science ☺



“I know how to pull a cart with help of 4 horses,

but I do not know how to pull a cart with help of 1024 chickens”

