




Chapter 4

Threads

Part **B**: Sections 4.2–4.4

Kernel-Level Threads (KLTs)



(b) Pure kernel-level

- **Thread management is done by the kernel**
 - There is no thread management code in the application level, simply an application programming interface (API) to the kernel thread facility
 - Windows is an example of this approach



Advantages of KLTs



- The kernel can simultaneously schedule multiple threads from the same process on multiple processors
- If one thread in a process is blocked, the kernel can schedule another thread of the same process
- Kernel routines themselves can be multithreaded, too, *for better OS-Performance*



Disadvantage of KLTs



❗ The transfer of control from one thread to another within the same process requires a mode switch to the kernel

Operation	User-Level Threads	Kernel-Level Threads	Processes
Null Fork	34	948	11,300
Signal Wait	37	441	1,840



Table 4.1


Thread and Process Operation Latencies (μ s)



Combined Approaches

- Thread creation is done completely in the user space, as is the bulk of the scheduling and synchronization of threads within an application

- Example: *Solaris*



(c) Combined



Threads:Processes	Description	Example Systems
1:1	Each thread of execution is a unique process with its own address space and resources.	Traditional UNIX implementations
M:1	A process defines an address space and dynamic resource ownership. Multiple threads may be created and executed within that process.	Windows NT, Solaris, Linux, OS/2, OS/390, MACH
1:M	A thread may migrate from one process environment to another. This allows a thread to be easily moved among distinct systems.	Ra (Clouds), Emerald
M:N	Combines attributes of M:1 and 1:M cases.	TRIX

Table 4.2
Relationship between Threads and Processes

Amdahl's Law !

- A formula with which we can *roughly estimate* the **speed-up** of a program (or *process*) when it is distributed over **several CPUs** in a concurrent mode of execution
- The formula is “*idealized*” because it **does NOT take into account the time-cost for thread-management** (scheduling, etc.) which the OS must invest to organize concurrency or parallelism. (*These time-costs can be quite large!*)
- For a program or process **P** the formula stipulates:
whereby:
 - $1 \leq \mathbf{s} \leq \mathbf{N}$ is the (*roughly estimated*) **speed-up factor**,
 - $1 \leq \mathbf{N} \leq \text{TLOC}$ is the **number of available CPUs** over which **P** may be distributed. *Note: more CPUs than Total Lines of Code would be useless, because the superfluous CPUs would stay idle!*
 - $0 \leq \mathbf{f} \leq 1$ is the (*estimated*) **fraction (percentage %)** of **P's total lines of code which is amenable to parallelization (concurrency)**, under the further assumption that there are no mutual-waiting-dependencies in this fraction. For accurate estimations the TLOC ought to be counted in atomic machine-code-instructions, not in “high level” software language.

$$\mathbf{s} = \frac{1}{(1 - \mathbf{f}) + \frac{\mathbf{f}}{\mathbf{N}}}$$

Estimation of f and s : a small example



- Consider this following “snippet” of three lines of code (L1–L3):

L1	$x := x + 1;$
L2	$y := y + 5;$
L3	$z := x + y;$



- The contents of L1 and L2 have nothing to do with each other: they are entirely independent of each other, and could thus be carried out separately and simultaneously by two different CPUs.



- L3, however, needs the results from L1 and L2 in order to be able to produce a value for **z**: hence, the computation of L3 *cannot happen simultaneously* with the computation of lines L1 and L2. L3 must be computed in a “sequential” mode and is thus “**not amenable**” for parallelisation in this small example.



- With only two out of these three lines being “amenable”, we get $f = 2/3 \approx 66\%$




- For these only 3 lines of code, **more than 3 CPUs would be useless**. In this example, **even N=3 would be useless, too**, because L3 is “not amenable”: L3 *could possibly be done by the same CPU that had already computed L1!* Thus, **2 CPUs will suffice (N=2)**.



- Hence, in this tiny example: $s \approx 1 / ((1 - 0.66) + (0.66 / 2)) \approx 1.49$

Speed-up *versus* Time saved

- Let **P** be a program (process) and **T** its **total run-time** which it would need on **1 CPU**.
- Let **s** be the **speed-up** value calculated for **P** by Amdahl's Law for **N CPUs**.
- **P**'s **shortened run-time T'** is thus (T / s) , whereby $T' < T$, if $s > 1$.
- The **time saved** is thus $S := T - T' = T - (T / s)$



- In our small *example from the previous slide*, with $s = 1.49$, the time saved would thus be $T - (T / 1.49)$
- Assume, *for example*, that $T=3$ (time-units for the *three lines of code*), then the **time saved** in this example would be $S = 3 - (3/1.49) \approx 0.98$ time-units

Amdahl's Law must be applied “wisely”!

- Consider this following “snippet”
of three lines of code (L1–L3):

L1	$x := x + 1;$
L2	$y := y + 5;$
L3	$z := x + y;$

- From the discussion of the previous slide we can learn that Amdahl's Law **must be applied “wisely”** in order to yield **meaningful and near-realistic speed-up estimations!**
- The *formula* of Amdahl's Law **can be easily mis-applied** because it *does not show the “hidden relation”* between the value **f** and the number of **N** CPUs which is “most appropriate” for this specific **f**!
- For example, we might *naively* guess: “*we have three lines of code here, thus let us set **N=3** and use 3 CPUs*”.
- Then we would *naively* calculate: $s \approx 1 / ((1-0.66)+(0.66/3)) \approx 1.79$, and we would *wrongly rejoice* about this apparently “good” speed-up, though in reality we cannot hope for more in this example than the **s=1.49** which we had derived on the previous slide. This is because the **3rd** CPU in this example **is idle** while the first two CPUs are still busy with the processing of lines L1 and L2!








Moreover:



- *In some cases it is not even possible to properly estimate the fraction-value f which is needed for calculating the approximate speed-up value s by means of Amdahl's Law!*
- In such cases, where theoretical analysis is not feasible, an experimental approach (with many repeated run-time experiments) is the only way of determining the actual speed-up value s .
- The considerations on the following slides will illustrate this dilemma.



Illustration of the Difficulties of “working with” Amdahl’s Law



```
Proc_def Difficult() // sequential version without concurrency
{
  int r := 0;
  int c := 0;


  for (i=0, i<1000, i++) do // loop with thousand repetitions
  {
    if (r==0) then { r := call Random_generator(); } // make any positive number

    if (r > 0) then { c := call Collatz(r); // c counts how often Collatz() was looping on r
                     r := 0; }


    if (c > 0) then { if ((c mod 2)==1) then { print("odd") } else { print("even") }
                     c := 0; }
  }
}
```

Attention: This example is NOT in the textbook! Please familiarize yourself with the Collatz algorithm.

Illustration of the Difficulties of “working with” Amdahl’s Law




```
Proc_def Difficult() // concurrent versions with threads: each running 1000 loops
{
  int r := 0;
  int c := 0;
```




```
Thread_def A() {
  for (i=0, i<1000, i++) do // thousand times
  { await (r==0) then { r := call Random_generator(); } }
```

Thread **A** depends on
Thread **B** via the shared
Process variable **r**



```
Thread_def B() {
  for (i=0, i<1000, i++) do // thousand times
  { await (r > 0) then { c := call Collatz(r);
                        r := 0; } } }
```

Thread **C** depends on
Thread **B** via the shared
Process variable **c**



```
Thread_def C() {
  for (i=0, i<1000, i++) do // thousand times
  { await (c > 0) then { if ((c mod 2) == 1) then { print("odd") } else { print("even") }
                        c := 0; } } }
```



```
}
```

Attention: This example is NOT in the textbook!

Illustration of the Difficulties of “working with” Amdahl’s Law

Proc_def **Difficult()** *// concurrent versions with threads: each running 1000 loops*
{
 int **r** := 0;
 int **c** := 0;

“constant” code section, the “size” of which can be easily estimated for Amdahl’s formula.

Thread_def **A()** {
 for (i=0, i<1000, i++) do
 { **await** (**r**==0) then { **r** := call **Random_generator()**; } }

“constant” code section, the “size” of which can be easily estimated for Amdahl’s formula.

Thread_def **B()** {
 for (i=0, i<1000, i++) do
 { **await** (**r** > 0) then { **c** := call **Collatz**(**r**);
 r := 0; } } }

The run-time “size” of this code section *cannot* be estimated, because for some inputs **r** the **Collatz** Algorithm returns a result **c** very quickly, whereas for some inputs **r** the **Collatz** algorithm takes very long time before it yields the output **c** for which **Thread C** is dependently waiting (and so the **r:=0** for which **A** waits).

Thread_def **C()** {
 for (i=0, i<1000, i++) do
 { **await** (**c** > 0) then { if ((**c** mod2)==1) then { **print**(“odd”) } else { **print**(“even”) }
 c := 0; } } }


“constant” code section, the “size” of which can be easily estimated for Amdahl’s formula.

}

Hence we cannot estimate the Speed-Up for **Difficult()** with Amdahl’s Law even if we know the number of CPUs

Attention: This example is NOT in the textbook!

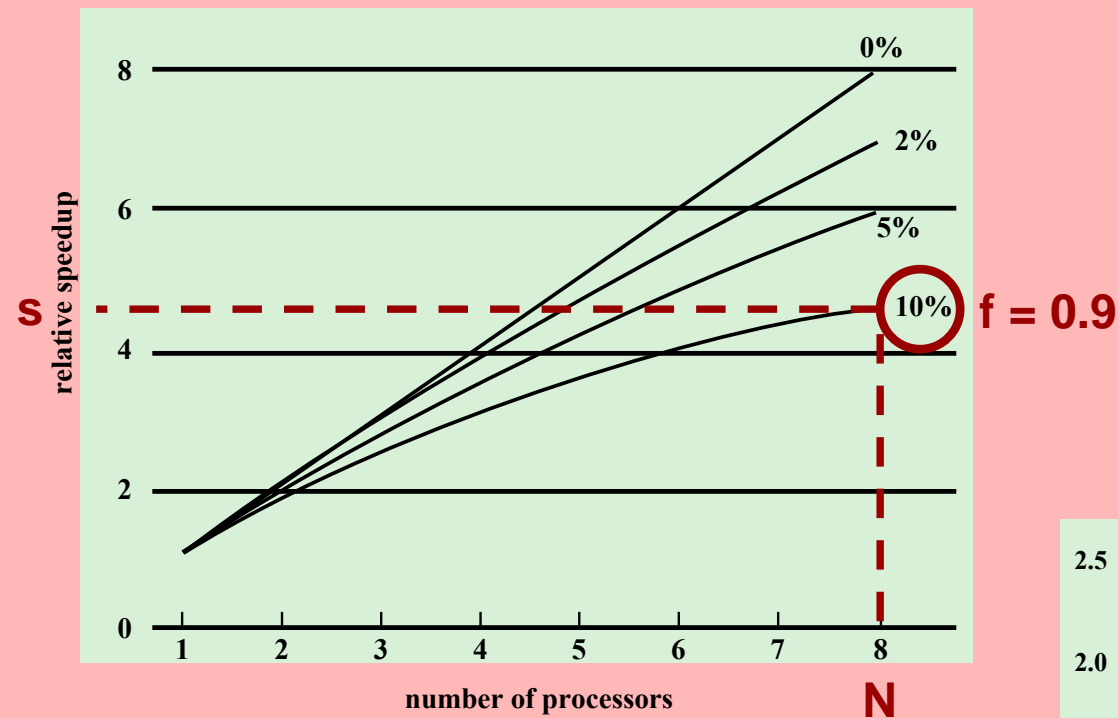
Thus we have learned:



If some code snippet appears inside a Loop, **its actual “size” depends on how often the Loop is actually repeated!** In many cases, however, this number of repetitions cannot be known beforehand




This makes it in some cases very difficult to estimate the “size” of some piece of Loop-Code, which *would* –however– be needed in order to give a reliable value to **f** in Amdahl’s Formula



(a) Speedup with 0%, 2%, 5%, and 10% sequential portions



My new Laptop
has 8 CPUs ☹️
Hehehe, my old
Laptop has only
4 CPUs 😊



(b) Speedup with overheads

Figure 4.7 Performance Effect of Multiple Cores




Figure 4.8 Scaling of Database Workloads on Multiple-Processor Hardware

Windows Process and Thread Management

- An **application** consists of one or more processes
- Each **process** provides the resources needed to execute a program
- A **thread** is the entity within a process that can be scheduled for execution
- A **job object** allows groups of process to be managed as a unit
- A **thread pool** is a collection of worker threads that efficiently execute asynchronous callbacks on behalf of the application
- A **fiber** is a unit of execution that must be manually scheduled by the application
- **User-mode scheduling (UMS)** is a lightweight mechanism that applications can use to schedule their own threads

Process and Thread *Objects*

Windows uses two types of process-related **objects**:



Processes

- An entity corresponding to a user job or application that owns resources

Threads

- A dispatchable unit of work that executes sequentially and is interruptible



→ A deeper understanding of “object-oriented” software will be taught in other courses of the Computer Science curriculum. **For now it suffices to regard an “object” as a “packaged combination” of suitably related Algorithms and Data Structures.**

A simple understanding of **“object”**

“Algorithms and their Data Structures in one Package”



- **Status** : *Ready*
- **Priority**: *High*
-
-

Algorithm:
Change_Status

Algorithm:
Change_Priority

Thread ID	A unique value that identifies a thread when it calls a server.
Thread context	The set of register values and other volatile data that defines the execution state of a thread.
Dynamic priority	The thread's execution priority at any given moment.
Base priority	The lower limit of the thread's dynamic priority.
Thread processor affinity	The set of processors on which the thread can run, which is a subset or all of the processor affinity of the thread's process.
Thread execution time	The cumulative amount of time a thread has executed in user mode and in kernel mode.
Alert status	A flag that indicates whether a waiting thread may execute an asynchronous procedure call.
Suspension count	The number of times the thread's execution has been suspended without being resumed.
Impersonation token	A temporary access token allowing a thread to perform operations on behalf of another process (used by subsystems).
Termination port	An interprocess communication channel to which the process manager sends a message when the thread terminates (used by subsystems).
Thread exit status	The reason for a thread's termination.

Table 4.4

Windows

Thread

Object's

Attributes



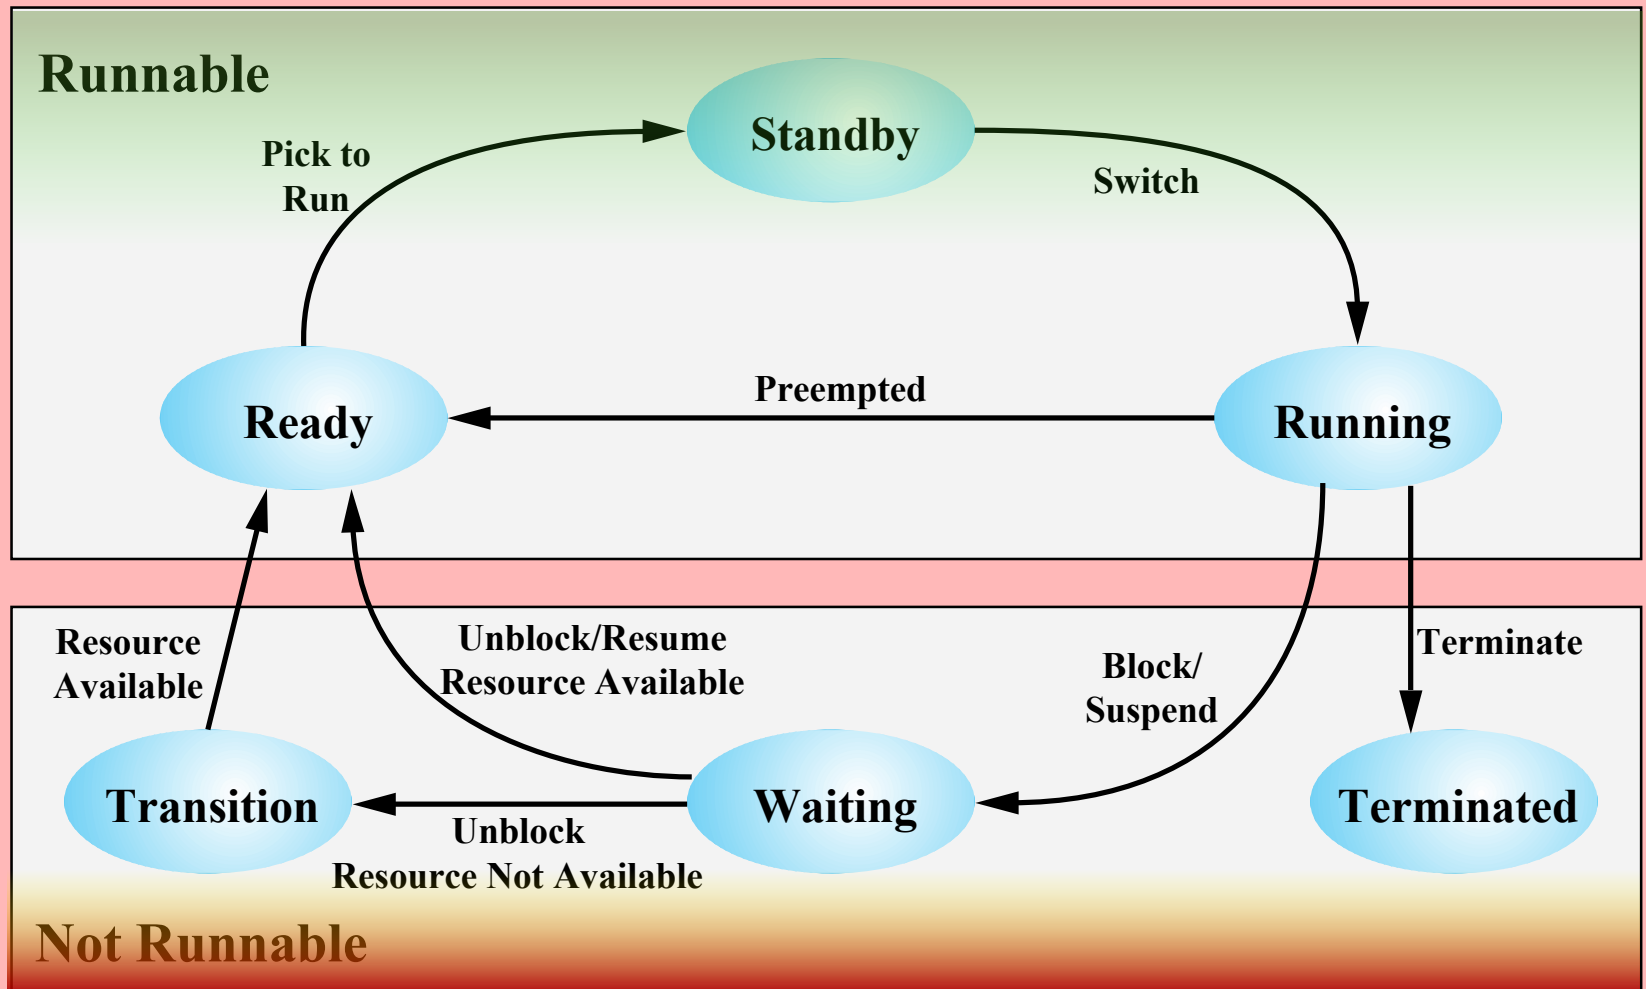
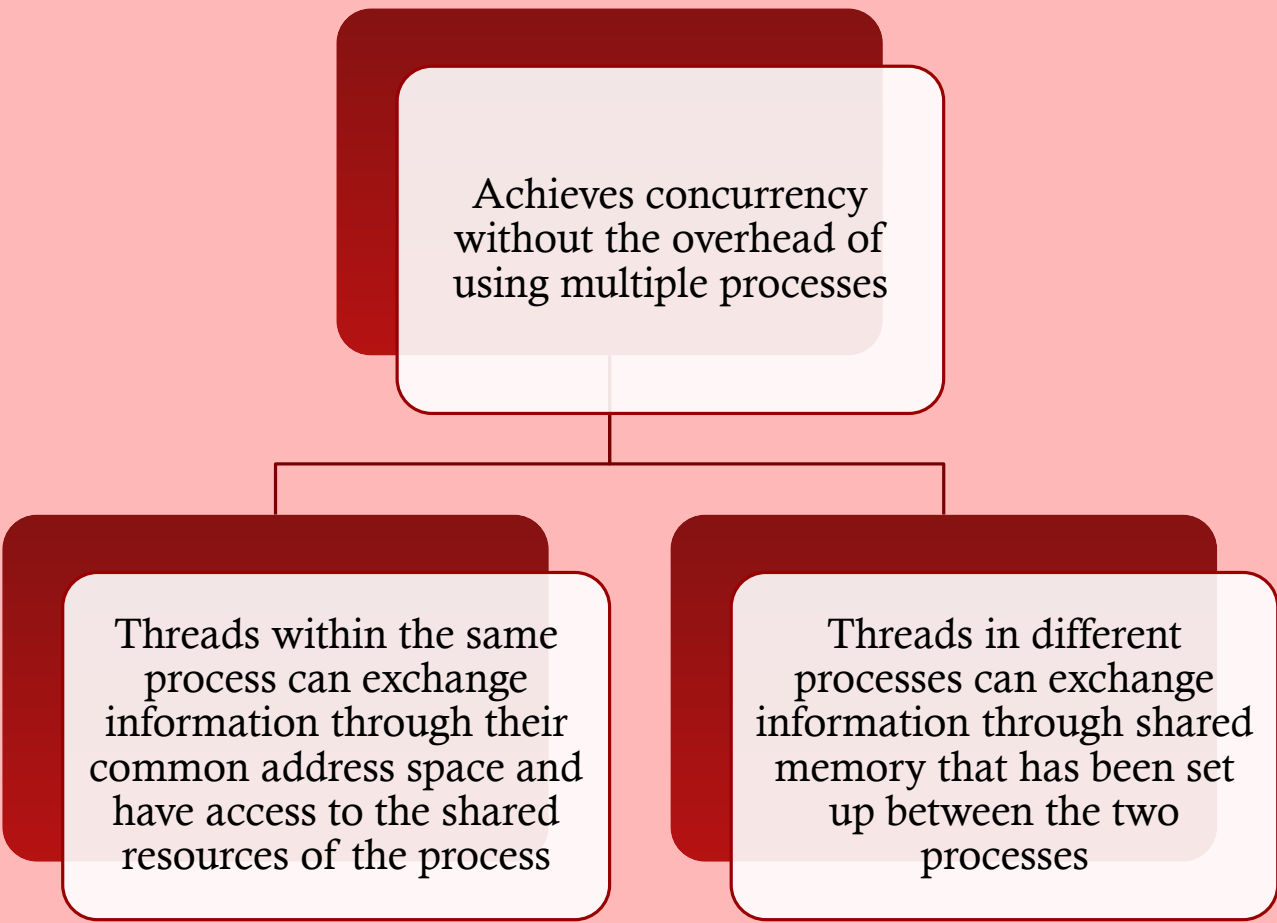


Figure 4.11 Windows Thread States



Multithreading



Achieves concurrency
without the overhead of
using multiple processes

Threads within the same
process can exchange
information through their
common address space and
have access to the shared
resources of the process

Threads in different
processes can exchange
information through shared
memory that has been set
up between the two
processes

