

# COS221 - L32 and L33 - Concurrency Control

Linda Marshall

22 and 24 May 2023

# Purpose of Concurrency Control

Focus is on the I (Isolation) and C (Consistency) of the ACID properties of a transaction.

- ▶ To ensure that the **isolation** property is maintained while allowing transactions to execute concurrently (outcome of concurrent transactions should appear as though they were executed in isolation).
- ▶ To preserve database **consistency** by ensuring that the schedules of executing transactions are serialisable.
- ▶ To resolve read-write and write-write conflicts among transactions.

# Concurrency Control Protocols

A Concurrency Control Protocol is a set of rules enforced by the DBMS to ensure serialisable schedules.

- ▶ Two important Concurrency Control Protocols are:
  - ▶ Two-phase locking (2PL)
  - ▶ Timestamps
- ▶ Other possible protocols include:
  - ▶ Multiversion Concurrency
  - ▶ Validation (Optimistic) Techniques
  - ▶ Multiple Granularity Locking

# Two-phase locking protocol

2PL locks data items. Locks describe the status of the item and therefore what operations can be applied to the item.

There are several types of locks that can be used for concurrency control, these include:

- ▶ Binary locks
- ▶ Shared/Exclusive (Read/Write) Locks

# Two-phase locking protocol - Binary Locks

Binary locks have two states, either the item is locked (1) or it is not (0) and each item in the database (X) has a distinct lock `LOCK(X)`.

- ▶ If the lock on item X is 1, then it cannot be accessed by the operation that requires that item.
- ▶ If the lock on item X is 0, then the lock can be obtained and the value of the lock is changed to 1

Two operations exist for binary locking:

- ▶ `lock_item(X)` - request to access item X
- ▶ `unlock_item(X)` - release item X

# Two-phase locking protocol - Binary Locks

Algorithms for the `lock_item(X)` and `unlock_item(X)` operations:

```
lock_item(X):  
B:  if  $\text{LOCK}(X) = 0$            (*item is unlocked*)  
    then  $\text{LOCK}(X) \leftarrow 1$    (*lock the item*)  
    else  
        begin  
        wait (until  $\text{LOCK}(X) = 0$   
            and the lock manager wakes up the transaction);  
        go to B  
        end;  
unlock_item(X):  
     $\text{LOCK}(X) \leftarrow 0;$       (* unlock the item *)  
    if any transactions are waiting  
        then wakeup one of the waiting transactions;
```

**Figure 21.1**  
Lock and unlock operations  
for binary locks.

# Two-phase locking protocol - Binary Locks

- ▶ A binary lock enforces **Mutual Exclusion** on the data item.
- ▶ The operations must be implemented as indivisible units (referred to as **critical sections**)
- ▶ The system maintains records for the items that are currently locked in a **lock table** which is organised as a hash file on the item name. Items not in the lock table are unlocked.
- ▶ The **lock manager subsystem** keeps track of and controls access to locks.

# Two-phase locking protocol - Binary Locks

For Binary Locks, every transaction must obey the following rules:

1. A transaction  $T$  must issue the operation `lock_item(X)` before any `read_item(X)` or `write_item(X)` operations are performed in  $T$ .
2. A transaction  $T$  must issue the operation `unlock_item(X)` after all `read_item(X)` and `write_item(X)` operations are completed in  $T$ .
3. A transaction  $T$  will not issue a `lock_item(X)` operation if it already holds the lock on item  $X$ .
4. A transaction  $T$  will not issue an `unlock_item(X)` operation unless it already holds the lock on item  $X$ .



# Two-phase locking protocol - Binary Locks

Binary locks are too restrictive.

- ▶ At most one transaction can hold a lock on a given item.
- ▶ If all transactions are only reading the item, then one transaction holding a lock on the item will result in the other transactions waiting until the transaction releases the lock.

Need a more granular locking system.

# Two-phase locking protocol - Shared/Exclusive Locks

Additional requirements for locking:

- ▶ Concurrent access to an item should be allowed for reading purposes. Reading an item by several transactions is not conflicting.
- ▶ A transaction writing to an item must have exclusive access to the item

Require a multimode lock, in this case the shared/exclusive (read/write) lock with three locking operations:

- ▶ `read_lock(X)`: transaction T requests a read (shared) lock on item X
- ▶ `write_lock(X)`: transaction T requests a write (exclusive) lock on item X
- ▶ `unlock(X)`: transaction T unlocks an item that it holds a lock on (shared or exclusive)

# Two-phase locking protocol - Shared/Exclusive Locks

Algorithms for the `read_lock(X)` and `write_lock(X)` operations:

**read\_lock(X):**

```
B:  if LOCK(X) = "unlocked"
      then begin LOCK(X) ← "read-locked";
           no_of_reads(X) ← 1
           end
    else if LOCK(X) = "read-locked"
      then no_of_reads(X) ← no_of_reads(X) + 1
    else begin
          wait (until LOCK(X) = "unlocked"
               and the lock manager wakes up the transaction);
        go to B
      end;
```

**write\_lock(X):**

```
B:  if LOCK(X) = "unlocked"
      then LOCK(X) ← "write-locked"
    else begin
          wait (until LOCK(X) = "unlocked"
               and the lock manager wakes up the transaction);
        go to B
      end;
```

# Two-phase locking protocol - Shared/Exclusive Locks

Algorithms for the `unlock(X)` operation:

```
unlock (X):  
  if LOCK(X) = "write-locked"  
    then begin LOCK(X)  $\leftarrow$  "unlocked";  
         wakeup one of the waiting transactions, if any  
    end  
  else if LOCK(X) = "read-locked"  
    then begin  
         no_of_reads(X)  $\leftarrow$  no_of_reads(X) - 1;  
         if no_of_reads(X) = 0  
           then begin LOCK(X) = "unlocked";  
                wakeup one of the waiting transactions, if any  
           end  
    end;
```

**Figure 21.2**

Locking and unlocking operations for two-mode (read/write, or shared/exclusive) locks.

## Two-phase locking protocol - Shared/Exclusive Locks

For Shared/Exclusive Locks, the system must enforce the following rules:

1. A transaction  $T$  must issue the operation  $\text{read\_lock}(X)$  or  $\text{write\_lock}(X)$  before any  $\text{read\_item}(X)$  operation is performed in  $T$ .
2. A transaction  $T$  must issue the operation  $\text{write\_lock}(X)$  before any  $\text{write\_item}(X)$  operation is performed in  $T$ .
3. A transaction  $T$  must issue the operation  $\text{unlock}(X)$  after all  $\text{read\_item}(X)$  and  $\text{write\_item}(X)$  operations are completed in  $T$ .
4. A transaction  $T$  will not issue a  $\text{read\_lock}(X)$  operation if it already holds a read (shared) lock or a write (exclusive) lock on item  $X$ .
5. A transaction  $T$  will not issue a  $\text{write\_lock}(X)$  operation if it already holds a read (shared) lock or write (exclusive) lock on item  $X$ .
6. A transaction  $T$  will not issue an  $\text{unlock}(X)$  operation unless it already holds a read (shared) lock or a write (exclusive) lock on item  $X$ .

# Two-phase locking protocol - Shared/Exclusive Locks

Rules 4 and 5 may be relaxed allowing lock conversions.

- ▶ Lock upgrade: existing read lock to write lock  
*if  $T_i$  holds a read-lock on  $X$ , and no other  $T_j$  holds a read lock on  $X$  ( $i \neq j$ ) then it is possible to convert (upgrade)  $read\_lock(X)$  to  $write\_lock(X)$  else force  $T_i$  to wait until all other transactions  $T_j$  that hold read locks on  $X$  release their locks*
- ▶ Lock downgrade: existing write lock to read lock  
*if  $T_i$  holds a write lock on  $X$  (this implies that no other transaction can have any lock on  $X$ ) then it is possible to convert (downgrade)  $write\_lock(X)$  to  $read\_lock(X)$*

# Two-phase locking protocol - Guaranteeing Serialisability

A transaction is said to follow the 2PL protocol, if all locking operations precede the first unlock operation in the transaction. The transaction is divided into the expanding phase, where locks are acquired and the shrinking phase where locks are relinquished.

If lock conversions are allowed then upgrading of locks must be done in the expanding phase and downgrading of locks in the shrinking phase.

# Two-phase locking protocol - Guaranteeing Serialisability

| $T_1$  | $T_2$  |
|--|--|
| <pre>read_lock(Y); read_item(Y); unlock(Y); write_lock(X); read_item(X); X := X + Y; write_item(X); unlock(X);</pre> | <pre>read_lock(X); read_item(X); unlock(X); write_lock(Y); read_item(Y); Y := X + Y; write_item(Y); unlock(Y);</pre> |

Does not follow the 2PL protocol

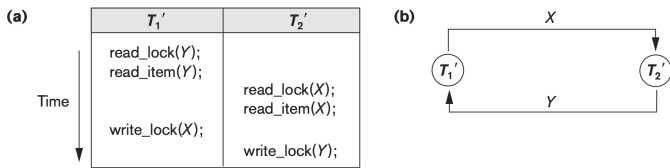
| $T_1'$   | $T_2'$   |
|--|--|
| <pre>read_lock(Y); read_item(Y); write_lock(X); unlock(Y); read_item(X); X := X + Y; write_item(X); unlock(X);</pre> | <pre>read_lock(X); read_item(X); write_lock(Y); unlock(X); read_item(Y); Y := X + Y; write_item(Y); unlock(Y);</pre> |

Follows the 2PL protocol  
but may produce a Deadlock



# Two-phase locking protocol - Deadlock and Starvation

A deadlock is a situation in which two or more transactions are waiting for one another to give up locks. But because the other transaction is also waiting (in a waiting queue), it will never release the lock.



**Figure 21.5**

Illustrating the deadlock problem. (a) A partial schedule of  $T_1'$  and  $T_2'$  that is in a state of deadlock. (b) A wait-for graph for the partial schedule in (a).

$T_1'$  is in the waiting queue for X, which is locked by  $T_2'$ , while  $T_2'$  is in the waiting queue for Y, which is locked by  $T_1'$ .

# Two-phase locking protocol - Variations

There are variations of the 2PL protocol:

- ▶ Conservative 2PL: A transaction must lock all its items before starting execution. If any item it needs is not available, it locks no items and tries again later. Conservative 2PL has no deadlocks since no lock requests are issued once transaction execution starts.
- ▶ Strict 2PL: All items that are writelocked by a transaction are not released (unlocked) until after the transaction commits. This is the most commonly used two-phase locking algorithm, and ensures strict schedules (for recoverability).
- ▶ Rigorous 2PL: All items that are writelocked or readlocked by a transaction are not released (unlocked) until after the transaction commits. Also guarantees strict schedules.

# Two-phase locking protocol - Deadlock and Starvation

## Deadlock prevention protocols

- A. Lock all items needed in advance - aka Conservative 2PL protocol - if any of the items cannot be obtained for locking, then none of the items are locked and the transaction waits and tries again later.
- B. Items in the database are ordered and transactions requiring the items lock them in that order.
- C. Handle a transaction that is in a possible deadlock situation by either:
  - ▶ waiting
  - ▶ aborting
  - ▶ preempting and aborting another transaction
- D. No Waiting (NW) and Cautious Waiting (CW) algorithms
- E. Deadlock detection

Options A and B limit concurrency. Option C makes use of **transaction timestamps**, while option D does not.

# Two-phase locking protocol - Deadlock and Starvation

## C. Transaction Timestamps

A transaction timestamp  $TS(T')$ , which is a unique identifier assigned to each transaction and is a monotonically increasing unique id given to each transaction based on their starting time.

If  $TS(T_1) < TS(T_2)$ , it means that  $T_1$  started before  $T_2$  ( $T_1$  older than  $T_2$ )

Two schemes that prevent deadlock are:

- ▶ Wait-die: An older transaction is allowed to wait for a younger transaction, whereas a younger transaction requesting an item held by an older transaction is aborted and restarted.
- ▶ Wound-wait: A younger transaction is allowed to wait for an older one, whereas an older transaction requesting an item held by a younger transaction preempts the younger transaction by aborting it.

# Two-phase locking protocol - Deadlock and Starvation

## D. Waiting algorithms

- ▶ No-waiting protocol: A transaction never waits; if  $T_i$  requests an item that is held by  $T_j$  in conflicting mode,  $T_i$  is aborted. Can result in needless transaction aborts because deadlock might have never occurred
- ▶ Cautious waiting protocol: If  $T_i$  requests an item that is held by  $T_j$  in conflicting mode, the system checks the status of  $T_j$ ; if  $T_j$  is not blocked, then  $T_i$  waits - if  $T_j$  is blocked, then  $T_i$  aborts. Reduces the number of needlessly aborted transactions

# Two-phase locking protocol - Deadlock and Starvation

## E. Deadlock Detection

- ▶ Deadlocks are allowed to happen. The system maintains a wait-for graph for detecting cycles. If a cycle exists, then one transaction involved in the cycle is selected (victim) and rolled-back (aborted).
- ▶ **Timeouts** can be used for deadlocks. If a transaction waits for a period longer than a predefined time, the system assumes the transaction is in deadlock and aborts irrespective of whether the transaction is in deadlock or not.
- ▶ **Starvation** occurs when a particular transaction consistently waits or gets restarted and never gets a chance to proceed further. A solution is to use a fair priority-based scheme that increases the priority for transaction the longer they wait e.g first-come-first-served queue.

# Two-phase locking protocol - Deadlock and Starvation

## Examples of Starvation

- ▶ In deadlock detection/resolution it is possible that the same transaction may consistently be selected as victim and rolled-back.
- ▶ In conservative 2PL, a transaction may never get started because all the items needed are never available at the same time.
- ▶ In Wound-Wait scheme a younger transaction may always be wounded (aborted) by a long running older transaction which may create starvation.

# Timestamps

Timestamps do not make use of locks and therefore deadlocks cannot occur.

Timestamps can be generated in several ways:

- ▶ Increment a counter each time its value is assigned to  $TS(T_i)$ . Counters have a finite value and therefore are periodically reset when there are no transactions running.
- ▶ Make use of the system clock and ensure that no two timestamp values are generated during the same tick of the clock.



# Timestamps - Timestamp Ordering Algorithm

Instead of locks, the system keeps track of two values for each data item X:

- ▶ **Read\_TS(X)**: The largest timestamp among all the timestamps of transactions that have successfully read item X
- ▶ **Write\_TS(X)**: The largest timestamp among all the timestamps of transactions that have successfully written X

When a transaction T requests to read or write an item X,  $TS(T)$  is compared with  $read\_TS(X)$  and  $write\_TS(X)$  to determine if request is out-of-order.

# Timestamps - Timestamp Ordering Algorithm

## Basic Timestamp Ordering

1. Transaction  $T$  requests a `write_item(X)` operation:
  - ▶ If  $\text{read\_TS}(X) > \text{TS}(T)$  or if  $\text{write\_TS}(X) > \text{TS}(T)$ , then a younger transaction has already read or written the data item so abort and roll-back  $T$  and reject the operation.
  - ▶ Otherwise, execute `write_item(X)` of  $T$  and set  $\text{write\_TS}(X)$  to  $\text{TS}(T)$ .
2. Transaction  $T$  requests a `read_item(X)` operation:
  - ▶ If  $\text{write\_TS}(X) > \text{TS}(T)$ , then a younger transaction has already written the data item  $X$  so abort and roll-back  $T$  and reject the operation.
  - ▶ Otherwise, execute `read_item(X)` of  $T$  and set  $\text{read\_TS}(X)$  to the larger of  $\text{TS}(T)$  and the current  $\text{read\_TS}(X)$ .

# Timestamps - Timestamp Ordering Algorithm

## Strict Timestamp Ordering

Schedules are both strict and serialisable.

Transaction  $T$  issues a `read_item(X)` or `write_item(X)` such that  $TS(T) > write\_TS(X)$  has its read or write operation delayed until the transaction  $T'$  that wrote the value of  $X$  (hence  $TS(T') = write\_TS(X)$ ) has committed or aborted.

# Timestamps - Timestamp Ordering Algorithm

## Thomas's Write Rule

This algorithm does not enforce conflict serialisability, but rejects fewer write operations as follows:

1. If  $\text{read\_TS}(X) > \text{TS}(T)$ , then abort and roll back  $T$  and reject the operation.
2. If  $\text{write\_TS}(X) > \text{TS}(T)$ , then do not execute the write operation but continue processing. This is because some transaction with timestamp greater than  $\text{TS}(T)$  - and hence after  $T$  in the timestamp ordering - has already written the value of  $X$ . Thus, we must ignore the  $\text{write\_item}(X)$  operation of  $T$  because it is already outdated and obsolete.
3. If neither the condition in 1. nor the condition in 2. occurs, then execute the  $\text{write\_item}(X)$  operation of  $T$  and set  $\text{write\_TS}(X)$  to  $\text{TS}(T)$ .

# Multiversion Concurrency

Multiversion concurrency control keeps multiple versions of the data item that have been written. When a read item is requested by a transaction the appropriate version is chosen to maintain serialisability.

Multiversion concurrency techniques in many cases make use of the concept of view serialisability rather than conflict serialisability.

A drawback of this technique is the added storage required to keep the multiple versions. However, some databases keep previous values of items making this technique a viable option.

# Multiversion Concurrency

Multiversion concurrency can be implemented using either:.

- ▶ Timestamps - The following two timestamps are kept for each version of  $X_i$ ,
  - ▶  $read\_TS(X_i)$ . The read timestamp of  $X_i$  is the largest of all the timestamps of transactions that have successfully read version  $X_i$ .
  - ▶  $write\_TS(X_i)$ . The write timestamp of  $X_i$  is the timestamp of the transaction that wrote the value of version  $X_i$ .
- ▶ 2PL protocol - Where three locking modes are used
  - ▶ read and write as before
  - ▶ certify

|     |       |      |       |
|-----|-------|------|-------|
| (a) |       | Read | Write |
|     | Read  | Yes  | No    |
|     | Write | No   | No    |

|     |         |      |       |         |
|-----|---------|------|-------|---------|
| (b) |         | Read | Write | Certify |
|     | Read    | Yes  | Yes   | No      |
|     | Write   | Yes  | No    | No      |
|     | Certify | No   | No    | No      |

**Figure 21.6**

Lock compatibility tables.  
(a) Lock compatibility table for read/write locking scheme.  
(b) Lock compatibility table for read/write/certify locking scheme.

# Validation (Optimistic) Techniques

The technique is optimistic because it assumes few conflicts will occur. The system therefore does not perform checks before reading and writing, but rather at the end of the transaction during the validation phase. The three phases occur in the following order:

1. Read phase - A transaction can read values of committed data items. However, writes are applied only to local copies (versions) of the data items.
2. Validation phase. Serialisability is checked by determining any conflicts with other concurrent transactions.
3. Write phase. On a successful validation, transaction updates are applied to the database on disk and become the committed versions of the data items; otherwise, transactions that fail the validation phase are restarted.

# Validation (Optimistic) Techniques

The validation phase for  $T_i$  checks that, for each transaction  $T_j$  that is either committed or is in its validation phase, one of the following conditions holds:

1.  $T_j$  completes its write phase before  $T_i$  starts its read phase.
2.  $T_i$  starts its write phase after  $T_j$  completes its write phase, and the  $\text{read\_set}(T_i)$  has no items in common with the  $\text{write\_set}(T_j)$
3. Both  $\text{read\_set}(T_i)$  and  $\text{write\_set}(T_i)$  have no items in common with the  $\text{write\_set}(T_j)$ , and  $T_j$  completes its read phase before  $T_i$  completes its write phase.



# Validation (Optimistic) Techniques

When validating  $T_i$ , the first condition is checked first for each transaction  $T_j$ , since 1. is the simplest condition to check.

If 1. is false for a particular  $T_j$ , then  
2. is checked and  
only if 2. is false is 3. checked.

If none of these conditions holds for any  $T_j$ , the validation fails and  $T_i$  is aborted.

# Multiple Granularity Locking


All database concurrency techniques assume that a database is a collection of named data items. A data item could be:

- ▶ A database record
- ▶ A field value of a database record
- ▶ A disk block
- ▶ A whole file
- ▶ The whole database

The choice of the granularity of a data item can influence the concurrency performance.

# Multiple Granularity Locking

By breaking up a database into a hierarchical structure as follows:



**Figure 21.7**

A granularity hierarchy for illustrating multiple granularity level locking.

a database can support multiple levels of granularity.

# Multiple Granularity Locking

To manage such hierarchy, in addition to read or shared (S) and write or exclusive (X) locking modes, three additional locking modes, called intention lock modes are defined:

- ▶ Intention-shared (IS): indicates that a shared lock(s) will be requested on some descendent nodes(s).
- ▶ Intention-exclusive (IX): indicates that an exclusive lock(s) will be requested on some descendent nodes(s).
- ▶ Shared-intention-exclusive (SIX): indicates that the current node is requested to be locked in shared mode but an exclusive lock(s) will be requested on some descendent node(s).

# Multiple Granularity Locking

|     | IS  | IX  | S   | SIX | X  |
|-----|-----|-----|-----|-----|----|
| IS  | Yes | Yes | Yes | Yes | No |
| IX  | Yes | Yes | No  | No  | No |
| S   | Yes | No  | Yes | No  | No |
| SIX | Yes | No  | No  | No  | No |
| X   | No  | No  | No  | No  | No |

**Figure 21.8**

Lock compatibility matrix for multiple granularity locking.

# Multiple Granularity Locking

The set of rules which must be followed for producing serialisable schedule are:

- ▶ The lock compatibility table must be adhered to.
- ▶ The root of the tree must be locked first, in any mode.
- ▶ A node N can be locked by a transaction T in S (or X) mode only if the parent node is already locked by T in either IS (or IX) mode.
- ▶ A node N can be locked by T in X, IX, or SIX mode only if the parent of N is already locked by T in either IX or SIX mode.
- ▶ T can lock a node only if it has not unlocked any node (to enforce 2PL policy).
- ▶ T can unlock a node, N, only if none of the children of N are currently locked by T.

# Multiple Granularity Locking

| $T_1$  | $T_2$   | $T_3$   |
|--|---|---|
| IX(db)<br>IX( $f_1$ )  | IX(db)  | IS(db)<br>IS( $f_1$ )<br>IS( $p_{11}$ )   |
| IX( $p_{11}$ )<br>X( $r_{111}$ )   | IX( $f_1$ )<br>X( $p_{12}$ )                        | S( $r_{111}$ )  |
| IX( $f_2$ )<br>IX( $p_{21}$ )<br>X( $p_{211}$ )                            |   |   |
| unlock( $r_{211}$ )<br>unlock( $p_{21}$ )<br>unlock( $f_2$ )               |   | S( $f_2$ )  |
|  | unlock( $p_{12}$ )<br>unlock( $f_1$ )<br>unlock(db) |   |
| unlock( $r_{111}$ )<br>unlock( $p_{11}$ )<br>unlock( $f_1$ )<br>unlock(db) |   | unlock( $r_{111}$ )<br>unlock( $p_{11}$ )<br>unlock( $f_1$ )<br>unlock( $f_2$ )<br>unlock(db) |

**Figure 21.9**

Lock operations to illustrate a serializable schedule.

# Other Concurrency techniques

Concurrency control can be applied to indexes. Two phase locking can be used in conjunction with B-Trees and  $B^+$ -Trees.

Insertion and Deletion of records may impact concurrency control resulting in a phantom record. A phantom record is a record that when being inserted satisfies the condition of another set of record accessed by another transaction. Phantom records can be dealt with by:

- ▶ index locking
- ▶ predicate locking

Interaction, e.g. by entering values on the screen and then submitting them may result in concurrency issues.