


The Stack and Functions

The Stack

- Up until now we have mostly ignored the stack.
- Instead we have relied on the data/bss segment.
- This has been sufficient so far, but this will change.
- We are now in a position where our programs will start being complex enough that we cannot rely entirely on fixed size allocations.



The stack in a Familiar Context

- Consider the following (intentionally inefficient) recursive C++ function. Assuming we don't know that

$$fib(n) = \frac{(1 + \sqrt{5})^n - (1 - \sqrt{5})^n}{2^n \sqrt{5}} \quad (1)$$

```
// assuming fib(1)=1
int fib(int n)
{
    if(n<2)
        return 1;
    int fibM1=fib(n-1);
    int fibM2=fib(n-2);
    return fibM1+fibM2;
}
```

- If N is too large our program will experience a stack overflow.

The Stack in a Familiar Context

- But why?

The Stack in a Familiar Context

- But why?
- There are two causes (though in a sense they are the same)
 - ▶ The more obvious one is each execution of the *fib* will push at least two integers onto the stack. One for each of *FibM1*, and *FibM2*. Now the space complexity of this naive fibonacci algorithm is $O(n)$. We can try and roughly calculate *fib*(10 000 000) we need $2 * 4\text{bytes} * 10\,000\,000 = 76.29\text{mb}$ just to store the max required number of *FibM1*s, and *FibM2*s.
 - ★ Most Linux distros have a stack between *8mb* and *16mb*

The Stack in a Familiar Context

- But why?
- There are two causes (though in a sense they are the same)
 - ▶ The more obvious one is each execution of the *fib* will push at least two integers onto the stack. One for each of *FibM1*, and *FibM2*. Now the space complexity of this naive fibonacci algorithm is $O(n)$. We can try and roughly calculate *fib*(10 000 000) we need $2 * 4\text{bytes} * 10\,000\,000 = 76.29\text{mb}$ just to store the max required number of *FibM1*s, and *FibM2*s.
 - ★ Most Linux distros have a stack between *8mb* and *16mb*
- What if we implement our recursive function to somehow not push any local variable data onto the stack?
 - ▶ Even if this was possible there is one stack item we would have to push.

The Stack in a Familiar Context


- But why?
- There are two causes (though in a sense they are the same)
 - ▶ The more obvious one is each execution of the *fib* will push at least two integers onto the stack. One for each of *FibM1*, and *FibM2*. Now the space complexity of this naive fibonacci algorithm is $O(n)$. We can try and roughly calculate *fib*(10 000 000) we need $2 * 4\text{bytes} * 10\,000\,000 = 76.29\text{mb}$ just to store the max required number of *FibM1*s, and *FibM2*s.
 - ★ Most Linux distros have a stack between *8mb* and *16mb*
- What if we implement our recursive function to somehow not push any local variable data onto the stack?
 - ▶ Even if this was possible there is one stack item we would have to push.
 - ▶ The return address.

The Stack

- Now that we have seen how it can break, lets us use the stack.
- The first point to recall is that the stack starts at :0x7fffffff. Unless there is some stack randomization in place.
- The second point is that the *rsp* register stores the stack pointer.
- If nothing has been added to the stack $rsp = 0x7fffffff$.
- We can interact with the stack in a number of ways, simplest of which are the `push` and `pop` instructions.

The Stack

- The `push` instruction decrements the `rsp` register and stores the value being pushed at this address
- The `pop` instruction places the value at the top of the stack into its operand and increments `rsp`
- With the x86-64 instructions you should push and pop 8 bytes at a time
 - ▶ It is also possible to push and pop 2 bytes (word) at a time
 - ▶ Direct 4 byte push and pops is not enabled in 64-bit mode



A Stack Example

Example with start of the stack at `0x7fffffff`

```
;A
mov  rax,74
push rax
;B
inc  rax
push rax
;C
inc  rax
push rax
;D
pop  rax
;E
pop  rax
;F
```

0x7fffffff0	x	74	74	74	74	74
0x7ffffffe8	x	x	75	75	75	x
0x7fffffe0	x	x	x	76	x	x
.	x	x	x	x	x	x
.	x	x	x	x	x	x
.	x	x	x	x	x	x
	A	B	C	D	E	F

In GDB

- You can use `x/1dg $rsp` to test the stack content.
- You can use `p/x $rsp` to test the stack pointer itself.
- Just remember that stack randomization effects the initial stack pointer.

The Stack

- Stack space is often reserved for local variables by subtracting the size needed from the stack pointer (`rsp`).
- Then an offset is used to refer to the variables.

```
sub rsp,      16 ;subtract 16 bytes
mov qword [rsp+8], 123 ;set our first qword variable to 123
mov qword [rsp],   24 ;set our second qword variable to 24
```

The stack

0x7fffffff				
0x7fffffff0	x		123	123
0x7fffffff08	x			24
0x7fffffff0e	x	x	x	x
.	x	x	x	x
.	x	x	x	x
.	x	x	x	x
	x	x	x	x
	sub rsp, 16 mov [rsp+8],123 mov [rsp], 24			

The Stack

- Remembering offsets for a large number of variables can become a burden. Instead use the `equ` pseudo-op.

```
first    equ 8
second   equ 0
mov qword [rsp+first],123
mov qword [rsp+second],24
```

The Stack

- Remembering offsets for a large number of variables can become a burden. Instead use the `equ` pseudo-op.

```
first    equ 8
second   equ 0
mov qword [rsp+first],123
mov qword [rsp+second],24
```

- How do we delete the variables after use?

The Stack

- Remembering offsets for a large number of variables can become a burden. Instead use the `equ` pseudo-op.

```
first    equ 8
second   equ 0
mov qword [rsp+first],123
mov qword [rsp+second],24
```

- How do we delete the variables after use?
 - ▶ We just move the stack pointer back.
`add rsp,16`

Functions

- Assuming we have loaded our parameters (will be explained shortly)
- You can call a function using

```
call    my_function
```

- `my_function` should be an appropriate address/label in the code segment
- The function's return value will be in `rax` or `xmm0`
- The effect of a function call is much like

```
push    next_instruction  
jmp     my_function
```

```
next_instruction:
```


The Return Instruction

- You can return to the location a function was called from using `ret`
- The effect of the return instruction (`ret`) is to pop an address off the stack and branch to it
- We could get much the same effect using

```
pop    rdi
jmp    rdi
```

Function Parameters

- How do we get our parameters to our function?
- On 32 bit Linux all parameters were pushed onto the stack

Function Parameters

- How do we get our parameters to our function?
- On 32 bit Linux all parameters were pushed onto the stack
- On x86-64 there are 8 more registers, so some parameters are passed in registers.

Function Parameters

- How do we get our parameters to our function?
- On 32 bit Linux all parameters were pushed onto the stack
- On x86-64 there are 8 more registers, so some parameters are passed in registers.
- Linux and Mac OS/X pass integer and address parameters 1 through 6 in `rdi`, `rsi`, `rdx`, `rcx`, `r8` and `r9`
 - ▶ The remaining integer and address parameters are pushed onto the stack
 - ▶ The first 8 floating point parameters are passed in registers `xmm0` - `xmm7`

Function Parameters

- How do we get our parameters to our function?
- On 32 bit Linux all parameters were pushed onto the stack
- On x86-64 there are 8 more registers, so some parameters are passed in registers.
- Linux and Mac OS/X pass integer and address parameters 1 through 6 in `rdi`, `rsi`, `rdx`, `rcx`, `r8` and `r9`
 - ▶ The remaining integer and address parameters are pushed onto the stack
 - ▶ The first 8 floating point parameters are passed in registers `xmm0` - `xmm7`
- Windows uses registers `rcx`, `rdx`, `r8` and `r9` for the first 4 integer and address parameters
 - ▶ The remaining integer and address parameters are pushed onto the stack. (there is also a set amount of stack padding)
 - ▶ Windows uses `xmm0` - `xmm3`

Function Parameters

- How do we get our parameters to our function?
- On 32 bit Linux all parameters were pushed onto the stack
- On x86-64 there are 8 more registers, so some parameters are passed in registers.
- Linux and Mac OS/X pass integer and address parameters 1 through 6 in `rdi`, `rsi`, `rdx`, `rcx`, `r8` and `r9`
 - ▶ The remaining integer and address parameters are pushed onto the stack
 - ▶ The first 8 floating point parameters are passed in registers `xmm0` - `xmm7`
- Windows uses registers `rcx`, `rdx`, `r8` and `r9` for the first 4 integer and address parameters
 - ▶ The remaining integer and address parameters are pushed onto the stack. (there is also a set amount of stack padding)
 - ▶ Windows uses `xmm0` - `xmm3`
- In all cases pushed parameters are pushed in reverse order

Function Parameters (2)

- Functions like `printf` having a variable number of parameters must place the number of floating point parameters in `rax`

Register	Usage	Preserved across function calls
<code>rax</code>	temporary register; with variable arguments passes information about the number of vector registers used; 1 st return register	No
<code>rbx</code>	callee-saved register; optionally used as base pointer	Yes
<code>rcx</code>	used to pass 4 th integer argument to functions	No
<code>rdx</code>	used to pass 3 rd argument to functions; 2 nd return register	No
<code>rsp</code>	stack pointer	Yes
<code>rbp</code>	callee-saved register; optionally used as frame pointer	Yes
<code>rsi</code>	used to pass 2 nd argument to functions	No
<code>rdi</code>	used to pass 1 st argument to functions	No
<code>r8</code>	used to pass 5 th argument to functions	No
<code>r9</code>	used to pass 6 th argument to functions	No
<code>r10</code>	temporary register, used for passing a function's static chain pointer	No
<code>r11</code>	temporary register	No
<code>r12-r15</code>	callee-saved registers	Yes

Simple Function

Simple function that returns the larger of two longs.

```
; long max(long a, long b)
```

```
max:
```

```
    mov rax, rdi    ; move parm1 to rax
```

```
    cmp rax, rsi    ; compare rax to parm2
```

```
    cmovl rax, rsi  ; if parm2 > rax then move parm 2 to rax
```

```
    ret
```


Simple Function

Calling the simple function

```
mov rdi, 123 ; load parm1  
mov rsi, 742 ; load parm2  
call max
```

Real World Functions

- Both Linux and Windows require the maintenance of the stack on 16 byte boundaries during the main part of functions.

Real World Functions

- Both Linux and Windows require the maintenance of the stack on 16 byte boundaries during the main part of functions.
 - ▶ The reason behind this requirement is to make it possible for local variables (on the stack) to be on 16 byte boundaries,
 - ▶ a requirement for some SSE(Streaming SIMD Extensions) and AVX(Advanced Vector Extension) instructions.

Real World Functions

- Both Linux and Windows require the maintenance of the stack on 16 byte boundaries during the main part of functions.
 - ▶ The reason behind this requirement is to make it possible for local variables (on the stack) to be on 16 byte boundaries,
 - ▶ a requirement for some SSE(Streaming SIMD Extensions) and AVX(Advanced Vector Extension) instructions.
- Conforming functions generally start with “push rbp” which re-establishes the 16 byte bounding temporarily **botched by the function call**
 - ▶ Remember, the call operation pushes a 8 byte value onto the stack (the return address)

Real World Functions

- Both Linux and Windows require the maintenance of the stack on 16 byte boundaries during the main part of functions.
 - ▶ The reason behind this requirement is to make it possible for local variables (on the stack) to be on 16 byte boundaries,
 - ▶ a requirement for some SSE(Streaming SIMD Extensions) and AVX(Advanced Vector Extension) instructions.
- Conforming functions generally start with “push rbp” which re-establishes the 16 byte bounding temporarily **botched by the function call**
 - ▶ Remember, the call operation pushes a 8 byte value onto the stack (the return address)
- Following that, conforming functions subtract multiples of 16 from rsp to allocate stack space or push pairs of 8 byte values
 - ▶ Even if this means over allocation

Hello world, at last

```
        section .data
msg:     db      "Hello World!",0x0a,0

        section .text
global   main
extern   printf

main:
    push    rbp
    mov     rbp, rsp      ; will explain shortly
    mov     rdi, msg      ; parameter 1 for printf
    mov     rax, 0        ; 0 floating point parameters
    call    printf
    mov     rax, 0        ; return 0
    mov     rsp, rbp      ; will explain shortly (NIB)
    pop     rbp
    ret
```

Stack frames

- Stack frames are used by the gdb debugger to trace backwards through the stack to inspect calls made in a process
- If we start and end each function like:

```
push    rbp
mov     rbp, rsp
```

```
...
mov     rsp, rbp
pop     rbp
ret
```

- We are in effect constructing a link list of all of the stack frames.
- All non-leaf functions must have the stack frame set up and destruction to conform to the ABI and be properly c/c++ compatible.

Stack frames

	In base function:	In Function L1:	In Function L2:	In Function L3:	In Function L4:
rbp	0x0	StackPointer(0)	StackPointer(1)	StackPointer(2)	StackPointer(3)
rsp	StackPointer(0)	StackPointer(1)	StackPointer(2)	StackPointer(3)	StackPointer(4)
0x7fffffff0	x	Ret Adr to Base	Ret Adr to Base	Ret Adr to Base	Ret Adr to Base
0x7fffffff8	x	0x0	0x0	0x0	0x0
0x7fffffffe0	x	x	Ret Adr to L1	Ret Adr to L1	Ret Adr to L1
0x7fffffffd8	x	x	StackPointer(0)	StackPointer(0)	StackPointer(0)
0x7fffffffd0	x	x	x	Ret Adr to L2	Ret Adr to L2
0x7fffffffc8	x	x	x	StackPointer(1)	StackPointer(1)
0x7fffffffc0	x	x	x	x	Ret Adr to L3
0x7fffffffb8	x	x	x	x	StackPointer(2)
0x7fffffffb0	x	x	x	x	x

*assuming no local variables are stored on the stack. If the local variables were stored where would they be?

Stack frames

- If you require space for local variables you simple need to subtract an amount from the stack pointer `rsp`.

Stack frames

- If you require space for local variables you simple need to subtract an amount from the stack pointer `rsp`.
 - ▶ The subtraction should always maintain the 16 byte boundary.
 - ▶ For example, say we wish to have a single quadword as a local variable.

Stack frames

- If you require space for local variables you simple need to subtract an amount from the stack pointer `rsp`.
 - ▶ The subtraction should always maintain the 16 byte boundary.
 - ▶ For example, say we wish to have a single quadword as a local variable.
 - ★ We only **have** to subtract 8 bytes.
 - ★ But we **should** maintain the 16 byte boundary.
- ```
push rbp
mov rbp, rsp
sub rsp,16
```

# Stack frames

- If you require space for local variables you simple need to subtract an amount from the stack pointer `rsp`.
  - ▶ The subtraction should always maintain the 16 byte boundary.
  - ▶ For example, say we wish to have a single quadword as a local variable.
    - ★ We only **have** to subtract 8 bytes.
    - ★ But we **should** maintain the 16 byte boundary.

```
push rbp
mov rbp, rsp
sub rsp,16
```

- if we allocate local variables we **must** use

```
mov rsp, rbp
pop rbp
ret
```

Just popping will not work.

# Stack frames

- If you prefer you can utilize

`leave`

`ret`

instead of

`mov rsp, rbp`

`pop rbp`

`ret`

## print max example

```
main:
 push rbp
 mov rbp, rsp
; print_max (100, 200);
 mov rdi, 100 ; first parameter
 mov rsi, 200 ; second parameter
 call print_max
 mov rax, 0 ; to return 0
 leave
 ret
```

## print max example

```
; void print_max (long a, long b)
; {
a equ 0
b equ 8
max equ 16
print_max:
 push rbp
 mov rbp, rsp
 sub rsp, 32 ; leave space for a, b and max
 mov [rsp+a], rdi ; save a
 mov [rsp+b], rsi ; save b
 mov [rsp+max], rdi ; max = a;
 cmp rsi, rdi ; if (b > max) max = b
 jng skip
 mov [rsp+max], rsi
```

## print max example

```
skip:
segment .data
 fmt db "max(%ld,%ld) = %ld",0xa,0
segment .text
 mov rdi, fmt ; address of format string
 mov rsi, [rsp+a] ; first %ld
 mov rdx, [rsp+b] ; second %ld
 mov rcx, [rsp+max] ; third %ld
 mov rax, 0 ; zero floating point param
 call printf
 leave
 ret
```



# Accessing Local Variables

- Can you see the potential issue with the way we are accessing our local variables?

# Accessing Local Variables

- Can you see the potential issue with the way we are accessing our local variables?
  - ▶ If anything is pushed onto the stack our offsets become incorrect.

# Accessing Local Variables

- Can you see the potential issue with the way we are accessing our local variables?
  - ▶ If anything is pushed onto the stack our offsets become incorrect.
- Solution?

# Accessing Local Variables

- Can you see the potential issue with the way we are accessing our local variables?
  - ▶ If anything is pushed onto the stack our offsets become incorrect.
- Solution?
  - ▶ Simply use `rbp` as the base address. Just remember that since `rbp` holds the value of `rsp` before the subtraction you need to use subtract from `rbp` and not add (think about what the offset values should be).
  - ▶ or just don't use push and pop other than for the stack frame setup and destruction

# Recursive Functions

- Recursive algorithms serve as a good example for why we need stack based storage.
- Often times we can get away with utilizing registers that are preserved across function calls.
- However consider the case where we have a recursive algorithm. On the first level we decide to use `r15` to store a value. But now on the second level `r15` is already in use...

# Recursive Functions

- Consider again the following recursive Fibonacci function

```
// assuming fib(1)=1
long fib(long n)
{
 if(n<2)
 return 1;
 return fib(n-1)+fib(n-2);
}
```

# Recursive Functions

```
• fib: push rbp
 mov rbp, rsp
 sub rsp, 16
N equ 0
nM1 equ 8
 mov rax, 1 ;base case return value
 cmp rdi, 2
 jl .end ;first parameter<2 (base case)
 dec rdi ;recall rdi is the first parameter (n)
 mov [rsp+N], rdi ;save N-1
 call fib
 mov [rsp+nM1], rax
 mov rdi, [rsp+N] ;load N-1
 dec rdi
 call fib
 add rax, [rsp+nM1]
 .end
 leave
 ret
```

# Function Implementation: Correct Practice

- If a function is a **non-leaf** function you **must** set up and destroy a stack frame.
- If you utilize a register in your function that should be preserved across function calls (like r15) you **must** restore it to its original value.
- Example

```
sub rsp, 16
mov [rsp], r15
mov [rsp+8], r14
....
....
mov r14, [rsp+8]
mov r15, [rsp]
add rsp, 16
```



# Function Implementation: Correct Practice

- or if you make use of no other stack based memory.

```
push r15
```

```
push r14
```

```
....
```

```
....
```

```
pop r14
```

```
pop r15
```