Chapter 9

Alternative Architectures

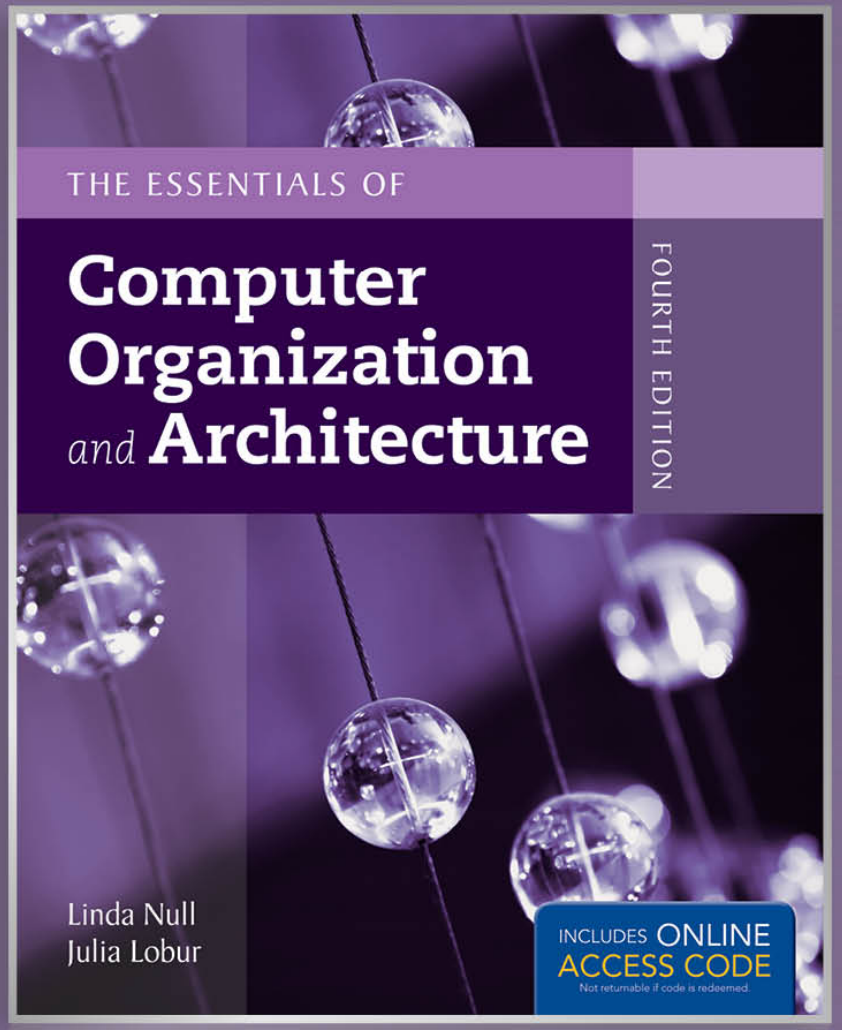

9.3 Flynn's Taxonomy

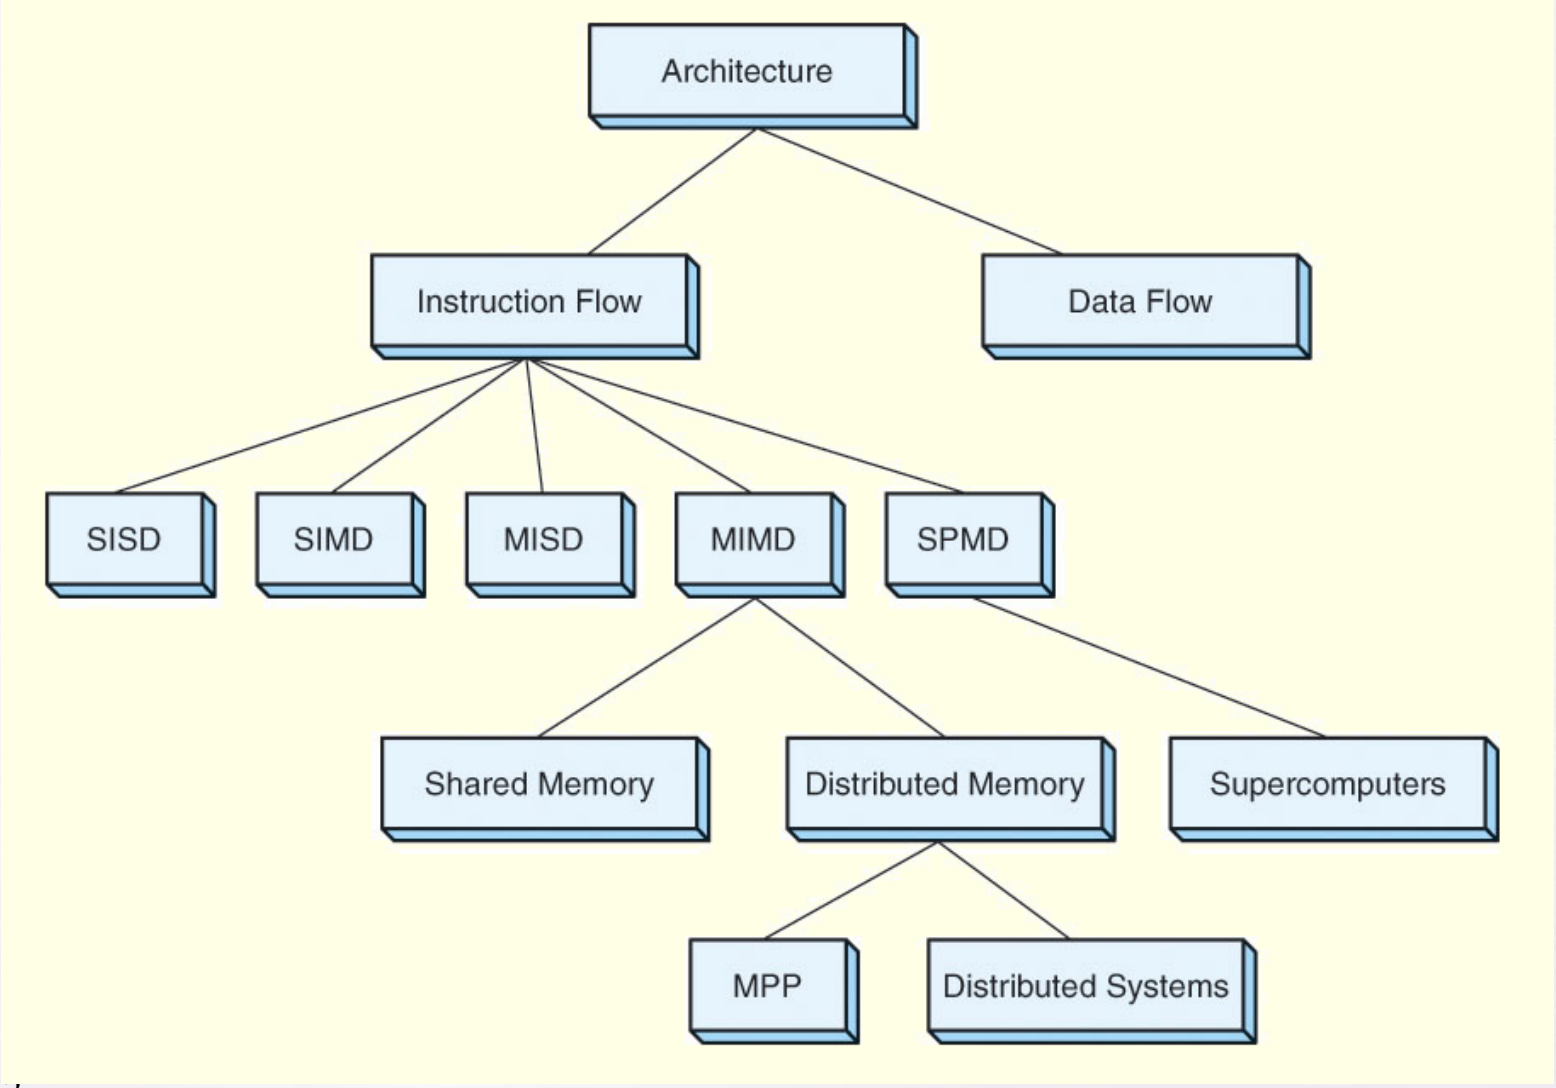

9.4 Parallel and Multiprocessor Architectures

- **Parallel processing** is capable of economically **increasing system performance**.

- The **limiting factor** is that no matter how well an algorithm is parallelized, there is always **some portion that must be done sequentially**.

- It is important to keep in mind that an ***n*-fold increase in processing** power does **not necessarily result in an *n* -fold increase in performance**.

9.4 Parallel and Multiprocessor Architectures

- **Pipelining divides** the **fetch-decode-execute** cycle **into single-cycle stages** that each carry out a small part of the process on a set of instructions.

- **Super pipelining** occurs when a **pipeline has substages that require less than half a clock cycle** to complete.

- The super pipeline is equipped with a **separate clock** running at a **frequency that is at least double** that of the main system clock.

9.4 Parallel and Multiprocessor Architectures

- **Superscalar architectures** include **multiple execution units** such as specialized integer and floating-point adders and multipliers.

- The superscalar **instruction fetch unit** can **simultaneously retrieve several instructions** from memory.

- The **hardware decoding unit determines which** of these **instructions can be executed in parallel and combines them**.

9.4 Parallel and Multiprocessor Architectures

- **Very long instruction word** (VLIW) architectures **use** a special **compiler** that **combines independent instructions into one long instruction** that is sent down the pipeline.

- One could argue that this is **better** than the hardware-based superscalar approach **because the compiler can better identify dependencies**.

- A counterargument is that **compilers good with static data but cannot have a view of the run time code**.

9.4 Parallel and Multiprocessor Architectures

- **Vector processors** are processors that **operate on entire vectors** or matrices at once.
- They are highly pipelined so that **arithmetic instructions can be overlapped**.
- Vector processors can be **categorized based on how operands are accessed**:
  - **Register-register** vector processors require all operands to be in registers.
  - **Memory-memory** vector processors allow operands to be sent from memory directly to the arithmetic units.

9.4 Parallel and Multiprocessor Architectures

- A **disadvantage** of **register-register** vector computers is that **large vectors must be divided into segments** to fit into the registers.

- **Memory-memory** vector computers have a **longer startup time until the pipeline becomes full**.

- Vector processors are **efficient because sets of values can be prefetched** since the processor knows it will have a continuous stream of data.

9.4 Parallel and Multiprocessor Architectures

- **MIMD** systems can **communicate through shared memory or** through an **interconnection network**.

- Interconnection networks are **classified** based on their **topology**, **routing** strategy, and **switching** technique.

- The **topology** is a **major determining factor** in the **overhead cost of message passing**.

9.4 Parallel and Multiprocessor Architectures

- **Interconnection networks** can be either **static or dynamic**.
- **Processor-to-memory connections** usually employ **dynamic** interconnections **based on switches**.
- **Processor-to-processor message-passing interconnections** are usually **static** and can employ any of several **different topologies**.

9.4 Processor-to-Processor Message-Passing Topologies

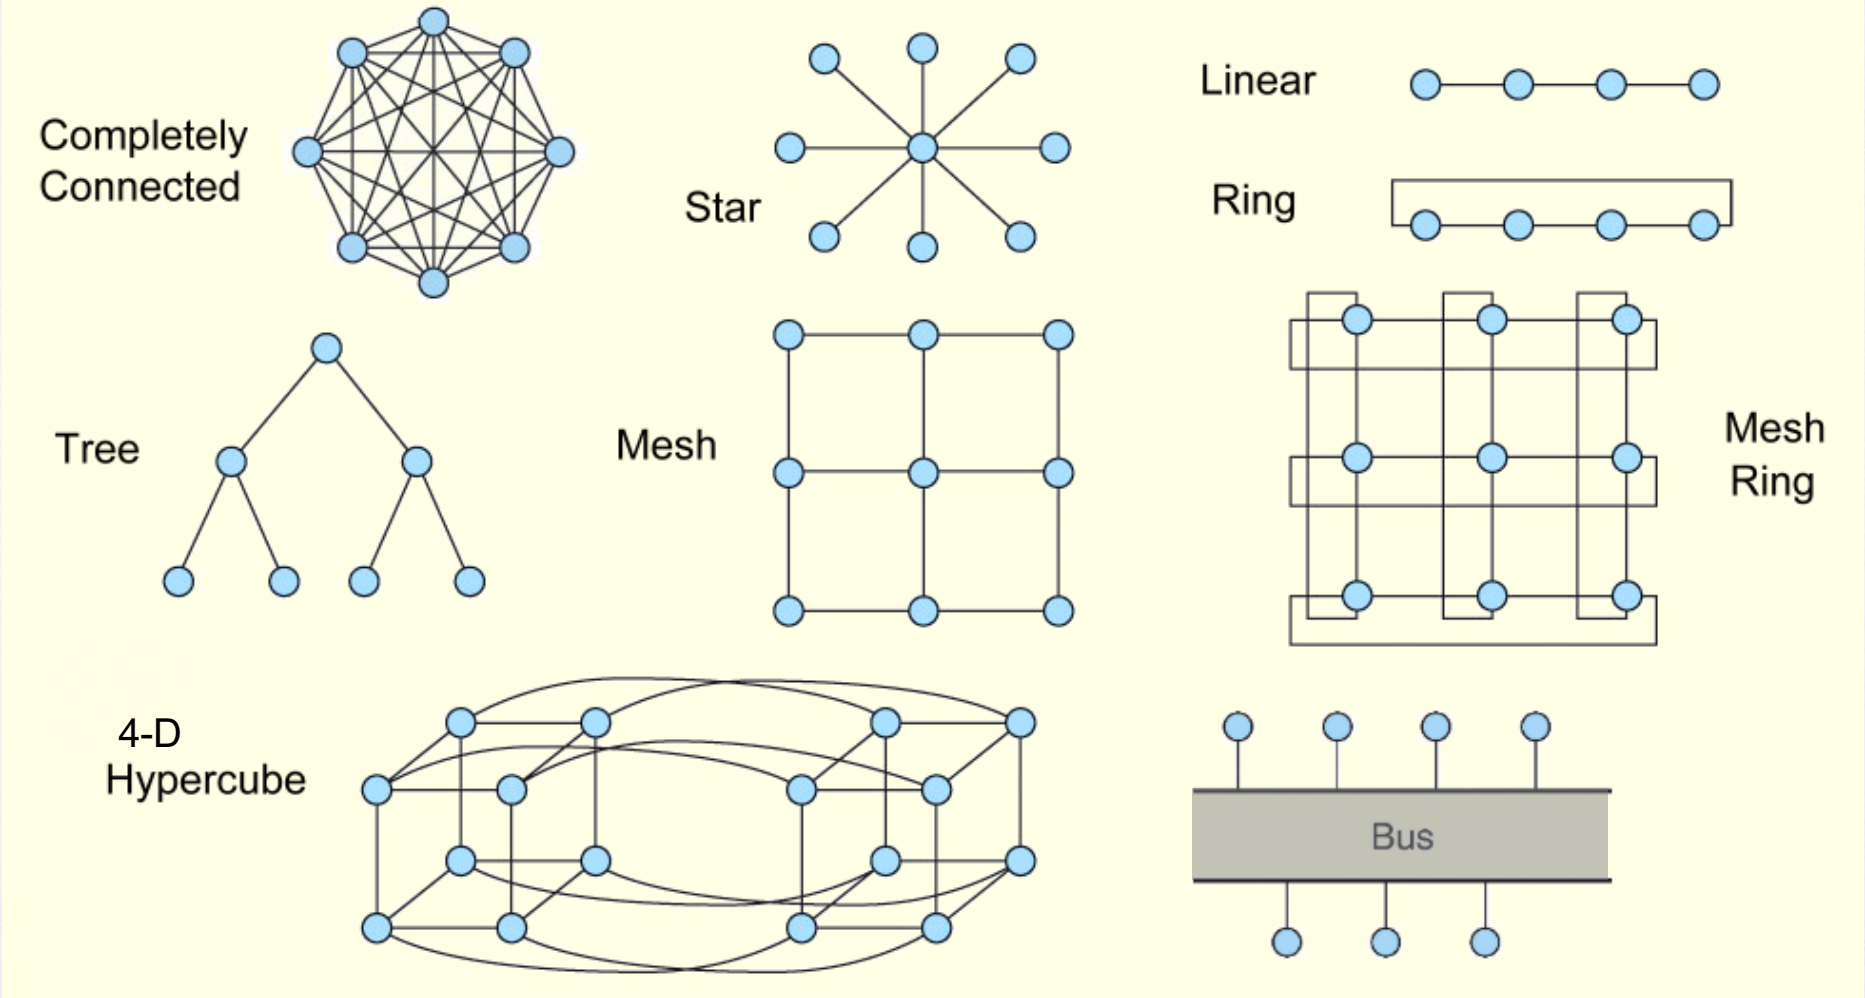

What is the max. length of the shortest path between arbitrary two nodes in these topologies?

9.4 Processor-to-Processor Message-Passing Topologies

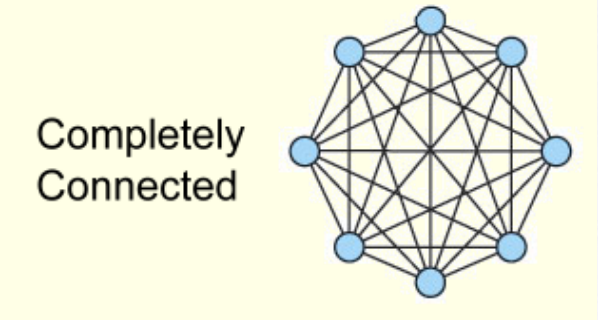

- Direct communication between any pair of nodes possible
- Durable network that isn't dependent on any one node
- Very secure due to 1-to-1 communication
- Suitable for small-sized networks

- Requires an extremely large amount of cable
- Takes a long time to set up
- Requires meticulous planning
- There is a limit to the number of cables each node can accommodate
- Difficult to add further nodes

9.4 Processor-to-Processor Message-Passing Topologies

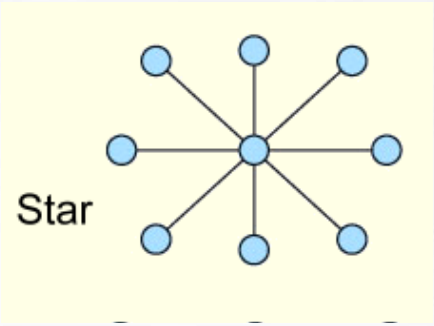

- At most two ‘hops’ for each pair of communicating nodes
- Easy to add and remove nodes
- Low cable usage
- Failure of an outer node will not affect other nodes in a network

- Requires specialist network hardware (hub node)
- A finite number of hub ports limits the network’s size
- Performance of transmission depends on the hub.
- Failure of the hub will stop any transmission.

9.4 Processor-to-Processor Message-Passing Topologies

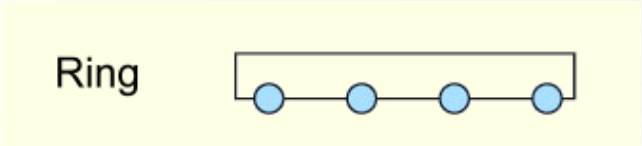

- Cheap to set up
- Easy to expand
- Dual ring option provides continuity through redundancy

- One faulty node or cable will bring the entire network down
- Performance declines with each additional node
- Reorganizing the network requires a full system shutdown

9.4 Processor-to-Processor Message-Passing Topologies

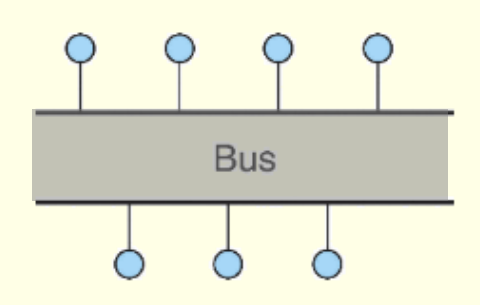

- Easy to connect a device and handle
- Takes less time to set up
- Best-suited for small networks.
- Easy to expand.

- If the bus cable fails, then the whole network will be down.
- Data can only travel in one direction at any point in time
- No communication can happen in parallel unless expensive parallel bus is used

9.4 Processor-to-Processor Message-Passing Topologies

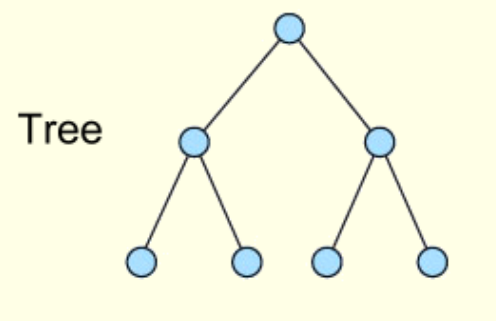

- Hybrid of bus and star topology
- Expandable
- Suitable for implementing divide-and-conquer approaches
- Network depends on the root node
- More cables needed in comparison to most other topologies

9.4 Parallel and Multiprocessor Architectures

- **Dynamic routing** is achieved **through switching** networks that consist of **crossbar** switches **or 2 × 2** switches.

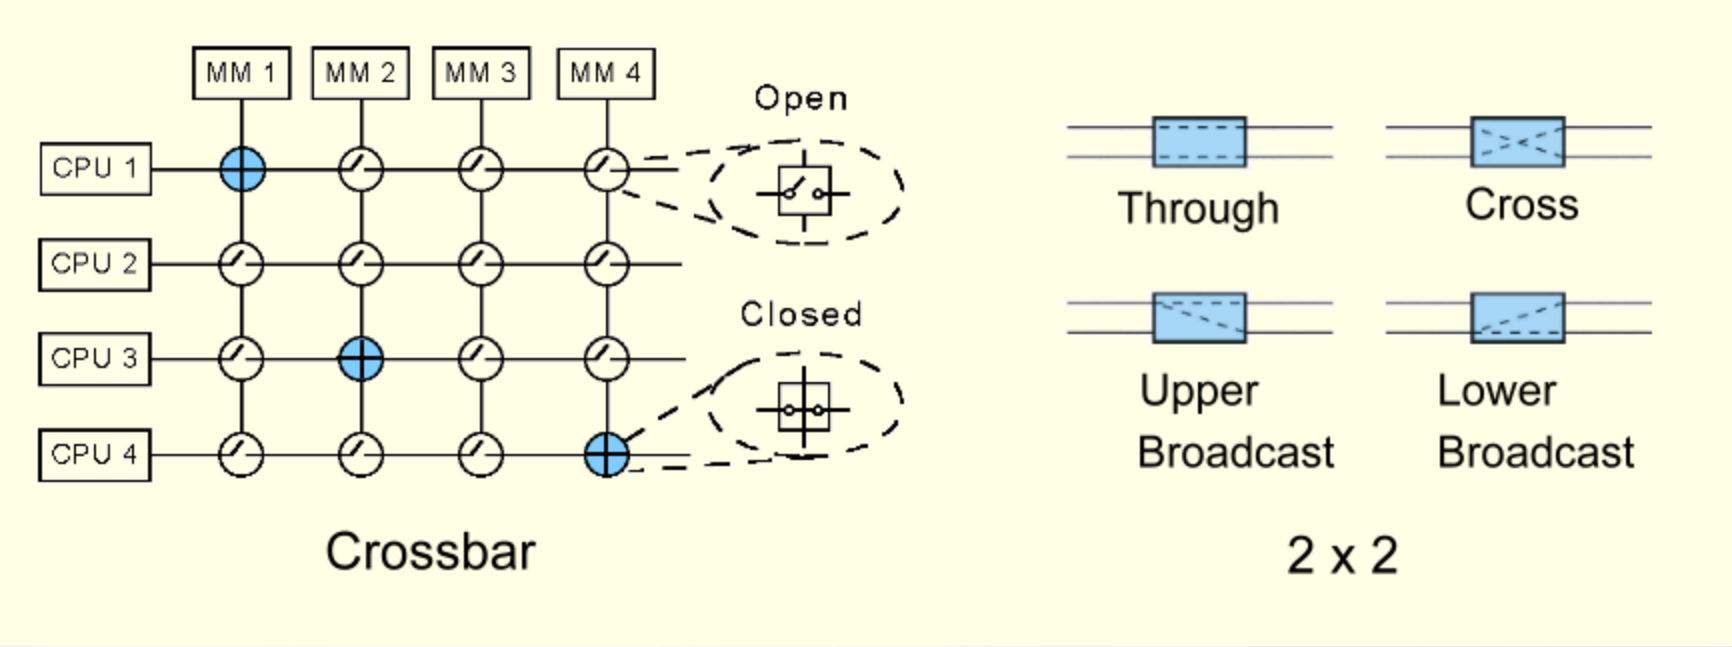

Quadratic number of switches required

9.4 Parallel and Multiprocessor Architectures

- **Multistage interconnection networks based on 2 x 2 switches**:

- $\log_2 n$ stages with n/2 switches per stage required

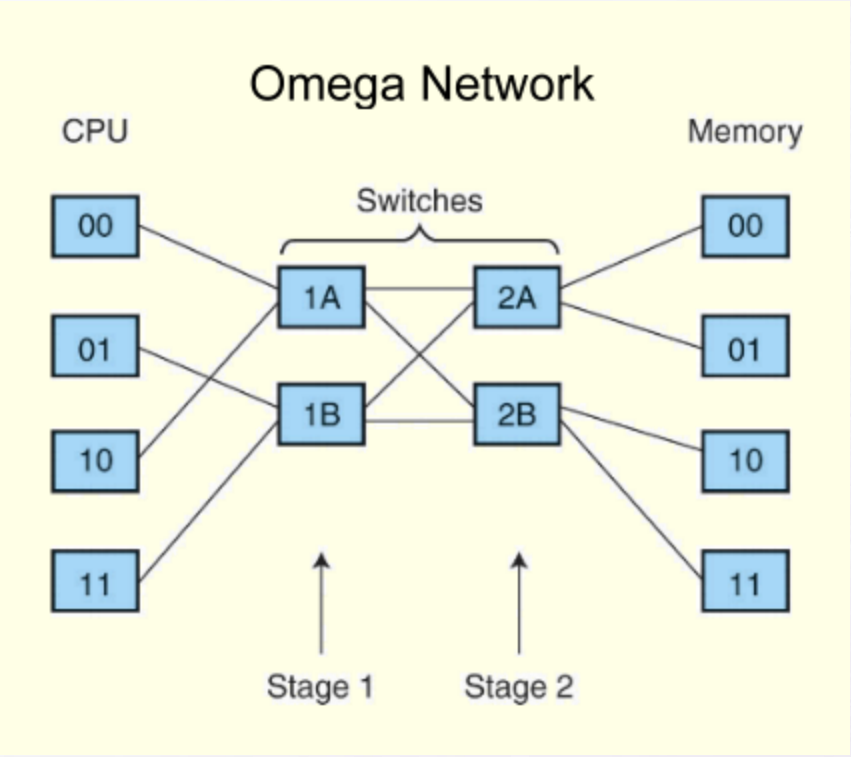

Chapter 9 Conclusion

- The common **distinctions between RISC and CISC systems** include RISC's short, fixed-length instructions. RISC ISAs are load-store architectures. These things permit RISC systems to be highly pipelined.

- **Flynn's Taxonomy** provides a way to **classify multiprocessor systems** based upon the number of processors and data streams.

Chapter 9 Conclusion

- **Massively parallel processors** have many processors, distributed memory, and computational elements communicate through a network. **Symmetric multiprocessors** have fewer processors and communicate through shared memory.

- Characteristics of **superscalar design** include **super pipelining**, and specialized instruction fetch and decoding units.

Chapter 9 Conclusion

- **Very long instruction word** (VLIW) architectures differ from superscalar architectures because the compiler, instead of a decoding unit, creates long instructions.
- **Vector computers** are highly-pipelined processors that operate on entire vectors or matrices at once.
- **MIMD systems** communicate through networks. The **network topology** often determines throughput.